

1 **18.337 FINAL PROJECT:**
2 **PERCEPTION-AWARE MULTIAGENT TRAJECTORY PLANNER**
3 **USING IMITATION LEARNING**

4 KOTA KONDO*

5 **Abstract.** Trajectory planning for unmanned aerial vehicles (UAVs) has been a focus of
6 extensive research; however, heavy computational requirements hinder their deployment in real-world
7 scenarios. One solution to this challenge is to use imitation learning planners that learn optimal tra-
8 jectories from existing planners and mimic their behavior. This approach offers the advantage of low
9 onboard computational requirements, making it more practical for real-world applications. While
10 single-agent trajectory planning has been extensively studied, multiagent planners have recently
11 gained popularity due to their broad range of applications, including package delivery. Multiagent
12 planners can be either centralized or decentralized, with the latter being more scalable and robust
13 to single-point failures.

14 Moreover, multiagent planners can be categorized as either asynchronous or synchronous, with
15 the former being more scalable than the latter. Perception-aware trajectory planning has become
16 increasingly popular among researchers due to its ability to gather information about the surrounding
17 environment and use it to plan trajectories. This approach is particularly useful for agents flying
18 in unfamiliar spaces. Although there have been numerous studies on perception-aware trajectory
19 planning for single agents, its use in multi-agent systems is still relatively uncommon.

20 To facilitate the training of perception-aware multiagent trajectory planners, we implemented
21 Message Passing Interface (MPI) on Julia, which is a standardized and portable message-passing
22 standard designed for parallel computing architectures. We conducted a performance comparison
23 that demonstrated MPI's advantage in parallelization.

24 Finally, we compared our imitation learning-based approach to optimization-based approaches
25 and found that our imitation learning approach had not been previously applied to decentralized,
26 asynchronous, perception-aware multiagent trajectory planners.

27 **Key words:** Julia [2], MPI, Imitation Learning, UAVs, Multiagent

28 **Codes:** <https://github.com/kotakondo/18337>

29 **1. Introduction.** In recent years, multiagent UAV trajectory planning has been
30 extensively studied [1, 4, 5, 7, 9, 12, 13, 15, 17, 18, 22, 25, 28, 31]. In real-world deploy-
31 ments of multiagent trajectory planning methods, it is crucial to deal with challenges
32 such as (1) detecting and avoiding collisions with **unknown obstacles**, (2) handling
33 **localization errors/uncertainties**, (3) achieving **scalability** to a large number of
34 agents, and (4) enabling **fast and efficient computation** for onboard replanning
35 and quick adaptation to dynamic environments. However, finding effective solutions
36 to these challenges remains an open question.

37 One approach to address challenges such as detecting and avoiding unknown
38 obstacles, even in the presence of localization errors and uncertainties, is to equip
39 each agent with a sensor, typically a camera, to perceive the surrounding environment.
40 This allows agents to gather real-time information about their surroundings, enabling
41 them to make informed decisions and take appropriate actions to avoid collisions and
42 navigate through dynamic environments. However, this sensor often has a limited field
43 of view (FOV), making the orientation of the UAV crucial when planning trajectories
44 through unknown space. Therefore, planners for flying with limited FOV sensors
45 generally need to be perception-aware to ensure that as many obstacles or other
46 UAVs as possible are kept within the FOV.

47 When scaling multiagent trajectory planners, it is important to note that, with
48 centralized planners, each agent needs to listen to a single entity that plans all the

*Department of Aeronautics and Astronautics, Massachusetts Institute of Technology, Cambridge,
MA (kkondo@mit.edu).

trajectories [12, 15]. While this approach simplifies planning, the central entity may act as a single point of failure, and the replanning abilities of the agent depend on their ability to communicate with the central entity. Decentralized planners greatly mitigate these issues, as each agent plans its own trajectory [1, 9, 18, 22, 25, 31]. Decentralized planners are therefore generally considered to be inherently more scalable and robust to failures.

Similarly, synchronous planners such as [4, 18, 27] require all agents to wait at a synchronization barrier until planning can be globally triggered, whereas asynchronous planning enables each agent to independently trigger the planning step without considering the planning status of other agents. Asynchronous methods are typically more scalable compared to synchronous methods [9, 22, 31]. Table 1 shows the categorization of scalability of these multiagent trajectory planning approaches.

Table 1. Multiagent Trajectory Planner Category

	Synchronous	Asynchronous
Centralized	Not Scalable	Not Possible
Decentralized	Somewhat Scalable	Most Scalable (our approach)

Many optimization-based approaches [9, 22, 25, 31] have been proposed for multiagent trajectory generation. However, these approaches often require substantial computational resources, posing challenges for deployments in dynamic environments that demand fast on-the-fly replanning. To mitigate this issue, researchers have explored imitation learning (IL)-based approaches [11, 20, 24], which offer the advantage of faster replanning while still achieving close-to-optimal trajectory generation.

To tackle the challenges of (1) **unknown objects detection and collision avoidance**, (2) **localization errors/uncertainties**, (3) **scalability**, and (4) **fast and efficient computation**, we propose an IL-based decentralized, asynchronous, perception-aware multiagent trajectory planner. Table 2 provides a comparison of the proposed approach with state-of-the-art approaches.

2. Trajectory Generation.

2.1. Expert — Optimization-based PA MA Planning. MADER [22] proposed an optimization-check-recheck scheme for decentralized, asynchronous multiagent planning. In this approach, an agent optimizes its trajectory while using received trajectories as optimization constraints. Next, the agent checks its trajectory against trajectories received in the optimization step and rechecks if it received any trajectory in the check step. To enhance robustness against communication delays, we proposed Robust MADER [9], which replaces the recheck step with a delay-check step. These frameworks allow fully decentralized asynchronous multiagent trajectory generation under real-world uncertainties and delays.

PANTHER [23] proposed a perception-aware trajectory planner for a single agent in dynamic environments, generating trajectories to avoid obstacles while keeping them in the sensor FOV. In [24], PANTHER* improved the original PANTHER with less conservatism and more optimal trajectory generation, but both were limited to tracking and avoiding only one obstacle at a time. To overcome this limitation, we

Table 2. State-of-the-art UAV Trajectory Planners

Method	Multiagent	Perception-aware
EGO-Swarm [31]		
DMPC [10]		
MADER [22]	Yes	No
decMPC [26]		
RMADER [9]		
Raptor [30]		
Time-opt [19]		
PANTHER [23]	No	Yes
PA-RHP [29]		
Deep-PANTHER [24]		
Proposed approach	Yes	Yes

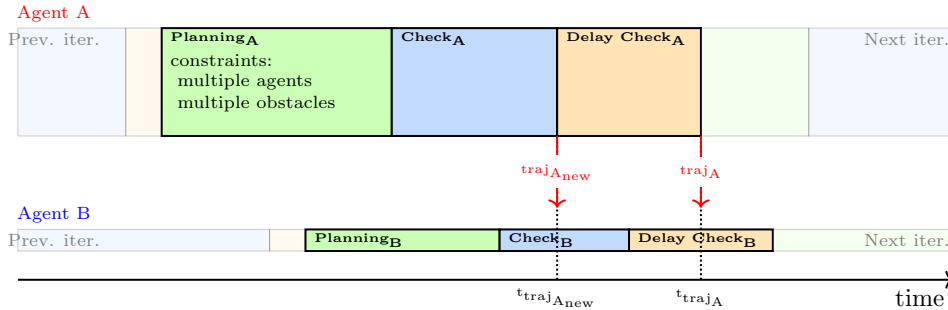


Fig. 1. Proposed trajectory optimization and deconfliction sequence: Our approach uses an imitation learning-based approach to generate trajectories for each agent, followed by a conflict detection and resolution step based on the Robust MADER framework. Each agent first generates a new trajectory in the planning step and then checks if there are any conflicts with the trajectories received from other agents. If no conflicts are detected, the agent publishes its new trajectory and begins checking for potential collisions in a delay check step. This delay check step is a sequence of checks over a period of time. Finally, if no conflicts are detected during the delay check, the agent commits to the new trajectory and publishes it. However, if conflicts are detected, the agent reverts to the trajectory from the previous iteration and discards the new trajectory. More details on the Robust MADER approach can be found in Section II of [9].

87 modified the optimization problem solved by PANTHER (see Appendix A) to enable
 88 tracking and avoidance of multiple obstacles, leading to a decentralized, asynchronous,
 89 perception-aware multi-agent trajectory planning system that incorporates this mod-
 90 ified optimization approach into the RMADER deconfliction framework. Fig. 1 illus-
 91 trates our approach’s trajectory deconfliction scheme, which is employed by both the
 92 expert and the student.

93 **2.2. Student —IL-based Approach.** Deep-PANTHER [24] used IL to train
 94 a neural network that generates a desired position trajectory, while relying on closed-
 95 form solutions to obtain the direction where the onboard sensor should be looking
 96 (e.g., yaw on a multirotor). This closed-form yaw solution generates yaw trajectories
 97 given position trajectories, reducing the output dimension of the learned policy. How-

98 ever, this approach is not scalable in multi-obstacle environments since the closed-form
 99 solution only generates yaw trajectories for a single given obstacle. To address this
 100 limitation, we designed our IL-based method using a multi-layer perceptron (MLP)
 101 that generates both position and yaw trajectories. To achieve this, we increased the
 102 size of the neural network to 4 fully connected layers, each with 1024 neurons, and
 103 trained it to imitate the optimal perception-aware trajectories.

104 Additionally, we added a Long Short-Term Memory (LSTM) [6] feature-extraction
 105 network to the MLP, inspired by the ground-robot motion planning approach [3]. This
 106 allowed the neural network to accept various numbers of obstacles and agents as input,
 107 whereas traditional feedforward neural networks can only handle a fixed number of
 108 obstacles. LSTM can take as many obstacles and agents as possible and generate a
 109 fixed length of the latent output, which we feed into the fully connected layers.

110 It is also worth noting that IL-based approaches are more scalable in practice than
 111 optimization-based approaches. As the number of agents and obstacles in the envi-
 112 ronment increases, optimization-based approaches need to include more constraints
 113 in the optimization, leading to significant computational requirements. On the other
 114 hand, IL-based approaches are able to handle larger-scale environments with little to
 115 no additional computational overhead with the use of LSTM.

116 In summary, we first fed the predicted trajectories of obstacles and received other
 117 agents’ trajectories to the LSTM, which outputs a fixed-size vector h . We then com-
 118 bine h with the agent’s own state and feed this into the fully connected layers. The
 119 architecture of the neural network is summarized in Fig. 2.

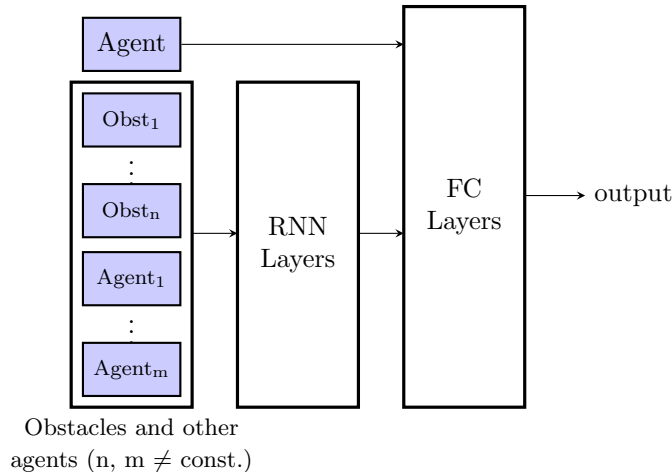


Fig. 2. Student Network Architecture

120 Table 3 shows the comparison of the state-of-the-art perception-aware trajectory
 121 planners. Our Expert approach is the first perception-aware multiagent trajectory
 122 planner that generates position and yaw coupled trajectory while tracking multiple
 123 obstacles, and our Student achieves much faster computation time, leveraging an
 124 IL-based planner.

¹Deep-PANTHER [24] generates only position trajectory, and yaw trajectory is generated by closed-form solution based on the position trajectory.

Table 3. State-of-the-art Perception-aware Obstacle Tracking Trajectory Planners

Method	Tracking Multi-obstacles	Multi-agents	Trajectory	Planning
[21]	No	No	Only Position	Optimization-based (slow & not scalable)
[14]	No	No	Position & Yaw	Optimization-based (slow & not scalable)
PANTHER / PANTHER* [23, 24]	No	No	Position & Yaw	Optimization-based (slow & not scalable)
Deep-PANTHER [24]	No	No	Only Position ¹	IL-based (faster & scalable)
Expert	Yes	Yes	Position & Yaw	Optimization-based (slow & not scalable)
Student (proposed)	Yes	Yes	Position & Yaw	IL-based (faster & scalable)

125 **2.3. Obstacle Sharing.** As shown in Fig. 3, each agent detects and tracks ob-
 126 stacles and shares their predicted trajectories with other agents. This obstacle-sharing
 127 architecture allows the agents to have a better understanding of the surrounding en-
 128 vironment as a team.

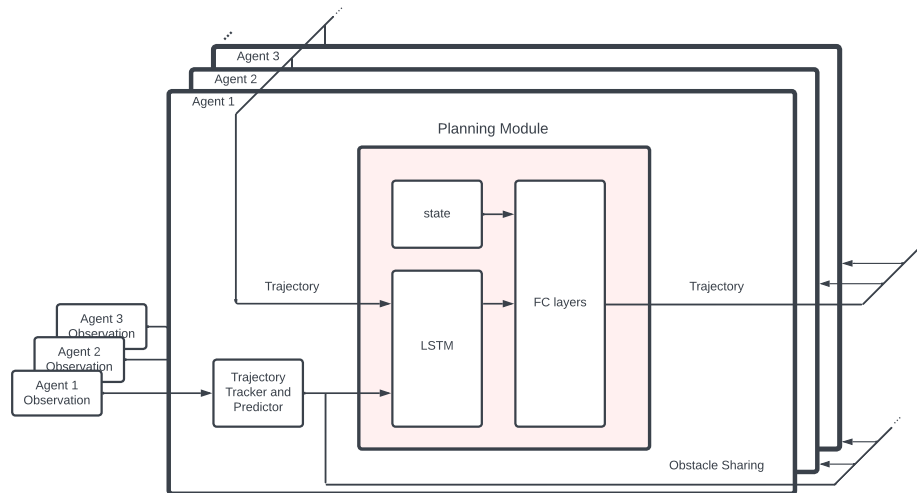


Fig. 3. Student Planning and Sharing Trajectory Architecture

129 3. Parallel Training.

130 **3.1. Training Setup.** We used the student-expert IL learning framework, where
 131 our expert approach provides demonstrations, and the student is trained so that
 132 its neural network can reproduce the provided demonstrations. The student was
 133 trained in an environment containing multiple dynamic obstacles following a random-

Table 4. Data Collection Time for 100 trajectories

		Data Collection Time [s]
Not Parallelized		320.4
MPI Parallelized	2 processors	181.6
	5 processors	86.8

134 ized trefoil-knot trajectory, with a randomized terminal goal. We first used Behavior
 135 Cloning (BC) to collect data and train the student, but its performance was sub-
 136 optimal compared to the expert. Therefore, we employed the Dataset-Aggregation
 137 algorithm (DAgger) [16] to refine the policy training using BC. The performance com-
 138 parison is given by Table 5, and Section 4 provides the detailed analysis. We used
 139 Adam [8] as an optimizer, and we normalized our observation and trajectory to make
 140 it easy for the neural network to learn. Additionally, we introduced a weighted loss
 141 function between position and yaw. During the training process, we found that it was
 142 more difficult to train the yaw trajectory than the position trajectory, and thus we
 143 weighted the yaw loss function. In our training, we set the weight α to 70. The total
 144 loss is defined as:

$$(3.1) \quad \mathcal{L}_{\text{total}} = \mathcal{L}_{\text{pos}} + \alpha \mathcal{L}_{\text{yaw}}$$

146 **3.2. Julia MPI performance comparison.** To accelerate the training pro-
 147 cess of perception-aware multiagent trajectory planners, we utilized the Julia MPI
 148 Package to implement parallelized training, which enables us to employ parallelized
 149 decentralized training, as depicted in Fig. 4. We performed a performance compar-
 150 ison and evaluated the training time using 1 (non-parallelized), 2, and 5 processors.
 151 Table 4 demonstrates that decentralized training completes much faster. Specifically,
 152 the two-processor and five-processor training completes 1.75 and 3.68 times faster
 153 than the non-parallelized training, respectively.

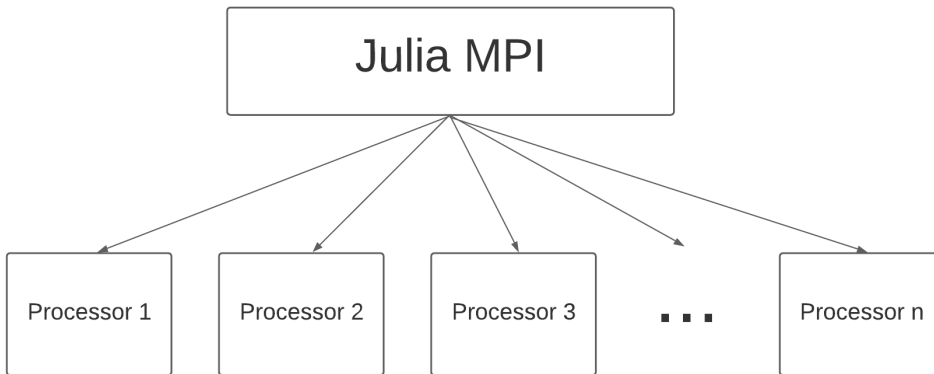


Fig. 4. Julia [2] MPI architecture

155 **4.1. Expert vs. Student in single-agent, single-obstacle environment.**
 156 Table 5 compares the average performance of the expert and the student in a sim-
 157 ulation environment with a single dynamic obstacle that follows a trefoil trajectory
 158 while the agent flies diagonally to avoid obstacles. The comparison is based on the
 159 average cost of trajectories and computation time, and the student with BC and DAg-
 160 ger achieves about 6000 times faster computation time with little performance loss.
 161 Fig. 6 shows the simulation environment.

Table 5. Expert vs. Student

	Avg. Cost	Computation Time [ms]
Expert	1317.0	5363.4
Student (BC)	2055.4	0.5634
Student (BC + DAgger)	1550.3	0.8978

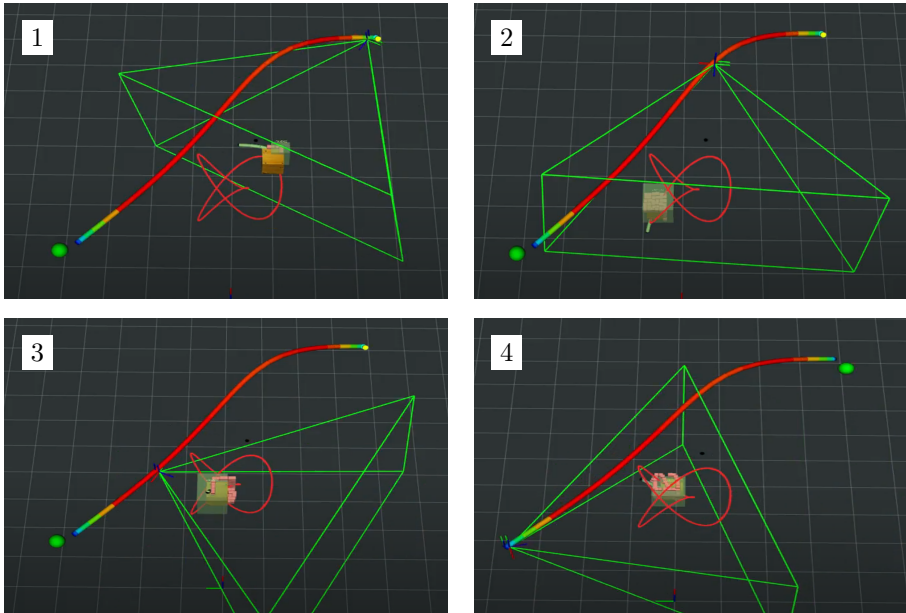


Fig. 5. Student single-agent, single-obstacle, simulation result: We made the imitation learning-based planner (student) fly around a trefoil-trajectory dynamic obstacle. The agent started at the top-right corner and was commanded to fly to the down-left.

162 **4.2. Multiagent and multi-obstacle benchmarking.** We also tested the ex-
 163 pert and the student in two different environments: one with one agent and two ob-
 164 stacles, and another with three agents and two obstacles. To conduct the experiment,
 165 we positioned the agents in a 3.0m radius circle and had them exchange positions

166 diagonally, as shown in Fig. 4. We set the maximum dynamic limits to 2.0 m/s,
 167 10.0 m/s², and 30.0 m/s³ for velocity, acceleration, and jerk, respectively.

168 We conducted all simulations on an Alienware Aurora r8 desktop running Ubuntu
 169 20.04, which is equipped with an Intel[®] Core[™] i9-9900K CPU clocked at 3.60 GHz
 170 with 16 cores and 62.6 GiB of RAM.

171 Table 6 and Fig. 7 compare the average performance of the expert and student
 172 in two different environments: (1) one agent with two obstacles, and (2) three agents
 173 with two obstacles. The metrics used to evaluate the performance are as follows:

- 174 1. Computation time: the time it takes to replan at each step.
- 175 2. Success rate: the rate at the agents successfully reach the goal without colli-
 176 sions.
- 177 3. Travel time: the time it takes for the agent to complete the position exchange.
- 178 4. FOV rate: the percentage of time that the agent keeps obstacles within its
 179 FOV when the agent is closer than its camera’s depth range.
- 180 5. Number of continuous FOV detection frames: the number of consecutive
 181 frames that an obstacle is kept within the FOV of the agent.
- 182 6. Dynamic constraints violation rate: the violation rate of the maximum ve-
 183 locity, acceleration, jerk, and yaw rate.

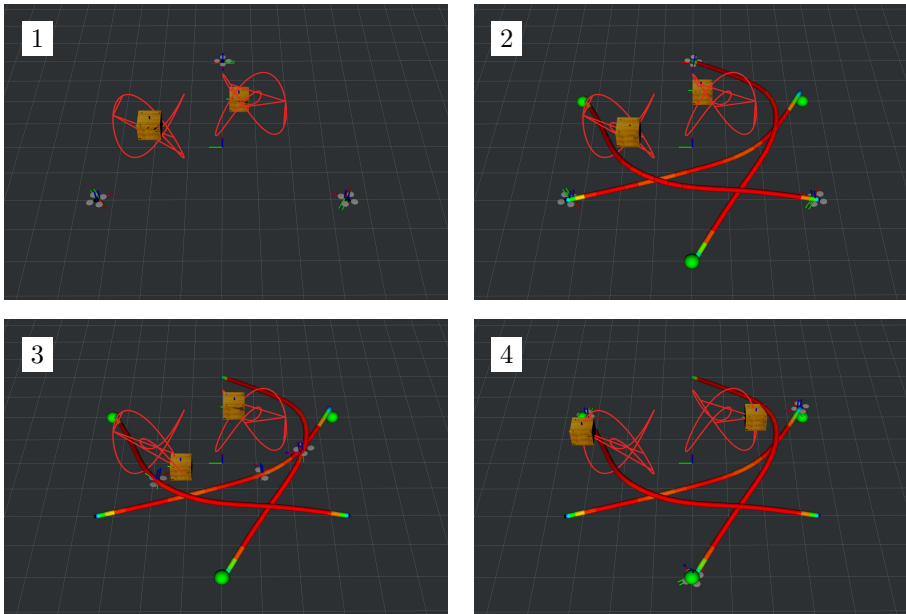


Fig. 6. Student mingle-agent, mingle-obstacle, simulation result: We made three imitation learning-based (student) agents fly around two dynamic obstacles. They started at the top-right corner and was commanded to fly to the down-left. For simplicity, we omitted FOV tripods visualization.

184 Both the expert and the studnet achieve successful position exchange with the
 185 two dynamic obstacles, with similar performance. However, the student significantly
 186 outperforms the expert in terms of computation time, completing the task in only
 187 57 ms compared to the much slower expert.

188 In the more complex environment with three agents and two obstacles, the expert
 189 and the student both achieve a high success rate, while the expert does not complete

Table 6. Benchmarking

Env.	Method	Avg. Compu. Time [ms]	Success Rate [%]	Avg. Travel Time [s]	FOV Rate [%]	Avg. of Max # Conti. FOV Detection Frames	Dyn. Constr. Violation Rate [%]
1 agent + 2 obst.	Expert	3456	100	7.9	29.0	19.8	0
	Student	57	100	4.5	28.0	31.0	10.3
3 agents + 2 obst.	Expert	6212	0	13.0	19.6	65.7	0.0
	Student	119	80	5.8	25.0	35.3	5.4

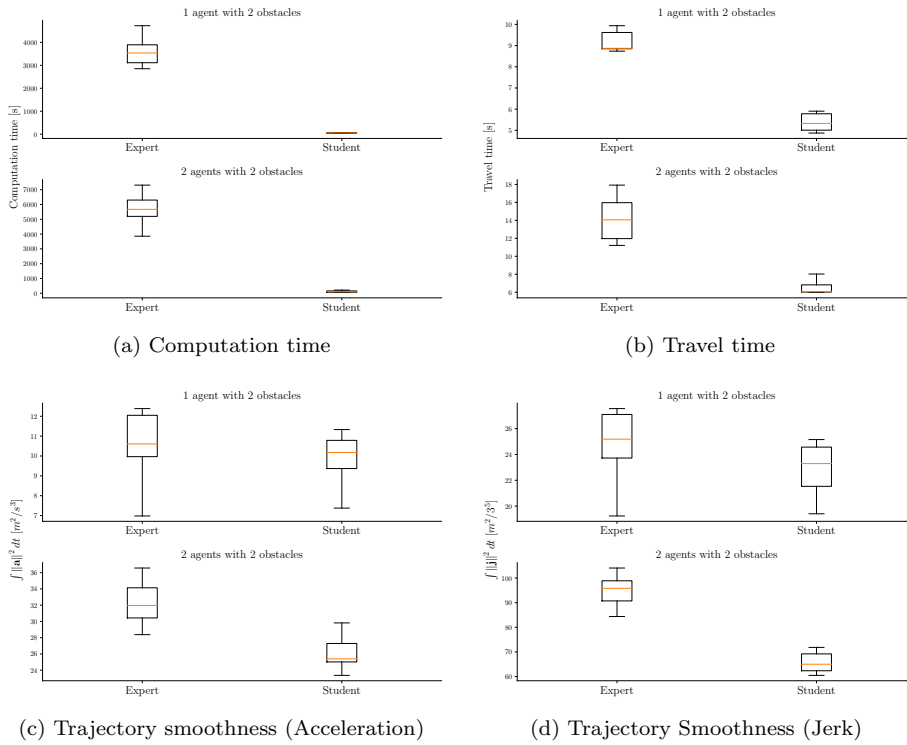


Fig. 7. Results of flight simulations. (a) The student’s computation time is much faster than that of the expert, and (b) the student’s travel time is also much shorter; this is mainly because of the faster computation time. (c-d) Since the student achieves faster replanning, it does not need to stop as the expert does, and that leads to smoother trajectory generation.

190 any position exchange. The reason for this is that when agents spend too much time
 191 optimizing their trajectory, the constraints used in the optimization become outdated
 192 by the time the optimization is complete, resulting in trajectory conflicts during the
 193 check and delay check steps. The expert suffers from long computation time, and as
 194 a result, almost none of the trajectories it generates pass the check and delay check
 195 steps, leading to a low success rate.

196 **5. Conclusions.** In conclusion, our work has addressed the critical issue of
 197 trajectory deconfliction in perception-aware, decentralized multiagent planning. We first
 198 presented an optimization-based perception-aware, decentralized, asynchronous multi-
 199 agent trajectory planner (expert) that enabled teams of agents to navigate uncertain
 200 environments while avoiding obstacles and deconflicting trajectories using perception
 201 information. Although these methods achieved state-of-the-art performance, they suf-
 202 fered from high computational costs, making it difficult for agents to replan at high
 203 rates.

204 To overcome this challenge, we presented a learning-based planner that was
 205 trained with imitation learning (IL). This approach is a computationally-efficient deep
 206 neural network and achieved a computation speedup of up to 6000 times faster than
 207 optimization-based approaches, while maintaining high performance. This speedup
 208 enables scalability to a large number of agents, making the student a promising ap-
 209 proach for large-scale swarm coordination. We used Julia MPI for parallelized training
 210 which led to 3.68 times faster training.

211 Moving forward, our future work will focus on larger-scale simulations and hard-
 212 ware flight experiments to demonstrate the scalability and performance of the student
 213 in complex environments with many agents and obstacles. Additionally, we will ex-
 214 plore how to integrate the student with other state-of-the-art perception systems,
 215 such as SLAM, to enable even more robust and accurate perception-aware multi-
 216 agent trajectory planning. Ultimately, our work has demonstrated the potential for
 217 learning-based approaches to address critical challenges in decentralized multiagent
 218 trajectory planning, and we believe that these approaches will play an essential role
 219 in enabling the deployment of multiagent systems in real-world applications.

220 Appendix A. Multi-obstacle Optimization Formulation.

221 To enable tracking multiple obstacles and agents we modified the FOV term given
 222 in Section 4 in [23] as the following.

$$223 \quad (A.1) \quad -\alpha_{FOV} \sum_i^n \left\{ \int_0^T (\text{inFOV}(\text{obstacle}_i))^3 dt \right\}$$

224 where α_{FOV} is the weight, n is the number of obstacles, T is the total time of the
 225 trajectory, $\text{inFOV}()$ returns a higher number when obstacle_i is in FOV.

226

REFERENCES

- 227 [1] S. BATRA, Z. HUANG, A. PETRENKO, T. KUMAR, A. MOLCHANOV, AND G. S. SUKHATME,
 228 *Decentralized control of quadrotor swarms with end-to-end deep reinforcement learning*, in
 229 Conference on Robot Learning, PMLR, 2022, pp. 576–586.
- 230 [2] J. BEZANSON, A. EDELMAN, S. KARPINSKI, AND V. B. SHAH, *Julia: A fresh approach to numer-
 231 ical computing*, SIAM review, 59 (2017), pp. 65–98, <https://doi.org/10.1137/141000671>.
- 232 [3] M. EVERETT, Y. F. CHEN, AND J. P. HOW, *Motion planning among dynamic, decision-making
 233 agents with deep reinforcement learning*, in 2018 IEEE/RSJ International Conference on
 234 Intelligent Robots and Systems (IROS), 2018, pp. 3052–3059, [https://doi.org/10.1109/
 235 IROS.2018.8593871](https://doi.org/10.1109/IROS.2018.8593871).
- 236 [4] R. FIROOZI, L. FERRANTI, X. ZHANG, S. NEJADNIK, AND F. BORRELLI, *A distributed
 237 multi-robot coordination algorithm for navigation in tight environments*, arXiv preprint
 238 arXiv:2006.11492, (2020).
- 239 [5] Y. GAO, Y. WANG, X. ZHONG, T. YANG, M. WANG, Z. XU, Y. WANG, Y. LIN, C. XU, AND
 240 F. GAO, *Meeting-merging-mission: A multi-robot coordinate framework for large-scale
 241 multi-robot coordination-limited exploration*, in 2022 IEEE/RSJ International Conference on In-
 242 telligent Robots and Systems (IROS), 2022, pp. 13700–13707, [https://doi.org/10.1109/
 243 IROS47612.2022.9981544](https://doi.org/10.1109/IROS47612.2022.9981544).

- 244 [6] S. HOCHREITER AND J. SCHMIDHUBER, *Long short-term memory*, Neural Computation, 9 (1997),
245 pp. 1735–1780.
- 246 [7] J. HOU, X. ZHOU, Z. GAN, AND F. GAO, *Enhanced decentralized autonomous aerial swarm*
247 *with group planning*, ArXiv, abs/2203.01069 (2022).
- 248 [8] D. P. KINGMA AND J. BA, *Adam: A method for stochastic optimization*, arXiv preprint
249 arXiv:1412.6980, (2014).
- 250 [9] K. KONDO, R. FIGUEROA, J. RACHED, J. TORDESILLAS, P. C. LUSK, AND J. P. HOW, *Rob-*
251 *ust mader: Decentralized multiagent trajectory planner robust to communication delay in*
252 *dynamic environments*, arXiv preprint arXiv:2303.06222, (2023).
- 253 [10] C. E. LUIS, M. VUKOSAVLJEV, AND A. P. SCHOELLIG, *Online trajectory generation with dis-*
254 *tributed model predictive control for multi-robot motion planning*, IEEE Robotics and Au-
255 tomation Letters, 5 (2020), pp. 604–611, <https://doi.org/10.1109/LRA.2020.2964159>.
- 256 [11] B. PARK AND H. OH, *Vision-based obstacle avoidance for uavs via imitation learning with*
257 *sequential neural networks*, International Journal of Aeronautical and Space Sciences, 21
258 (2020), pp. 768 – 779.
- 259 [12] J. PARK, J. KIM, I. JANG, AND H. J. KIM, *Efficient Multi-Agent Trajectory Planning with*
260 *Feasibility Guarantee using Relative Bernstein Polynomial*, in 2020 IEEE International
261 Conference on Robotics and Automation (ICRA), May 2020, pp. 434–440, <https://doi.org/10.1109/ICRA40945.2020.9197162>. ISSN: 2577-087X.
- 262 [13] P. PENG, W. DONG, G. CHEN, AND X. ZHU, *Obstacle avoidance of resilient uav swarm forma-*
263 *tion with active sensing system in the dense environment*, arXiv preprint arXiv:2202.13381,
264 (2022).
- 265 [14] B. PENIN, R. SPICA, P. R. GIORDANO, AND F. CHAUMETTE, *Vision-based minimum-time tra-*
266 *jectory generation for a quadrotor uav*, in 2017 IEEE/RSJ International Conference on
267 Intelligent Robots and Systems (IROS), 2017, pp. 6199–6206, <https://doi.org/10.1109/IROS.2017.8206522>.
- 270 [15] D. R. ROBINSON, R. T. MAR, K. ESTABRIDIS, AND G. HEWER, *An Efficient Algorithm for*
271 *Optimal Trajectory Generation for Heterogeneous Multi-Agent Systems in Non-Convex*
272 *Environments*, IEEE Robotics and Automation Letters, 3 (2018), pp. 1215–1222, <https://doi.org/10.1109/LRA.2018.2794582>. Conference Name: IEEE Robotics and Automation
273 Letters.
- 274 [16] S. ROSS, G. GORDON, AND D. BAGNELL, *A reduction of imitation learning and structured pre-*
275 *dition to no-regret online learning*, in Proceedings of the fourteenth international confer-
276 ence on artificial intelligence and statistics, JMLR Workshop and Conference Proceedings,
277 2011, pp. 627–635.
- 279 [17] G. RYOU, E. TAL, AND S. KARAMAN, *Cooperative Multi-Agent Trajectory Generation with*
280 *Modular Bayesian Optimization*, in Robotics: Science and Systems XVIII, Robotics: Science
281 and Systems Foundation, June 2022, <https://doi.org/10.15607/RSS.2022.XVIII.060>,
282 <http://www.roboticsproceedings.org/rss18/p060.pdf> (accessed 2022-07-08).
- 283 [18] B. SABETGHADAM, R. CUNHA, AND A. PASCOAL, *A distributed algorithm for real-time multi-*
284 *drone collision-free trajectory replanning*, Sensors, 22 (2022), <https://doi.org/10.3390/s22051855>, <https://www.mdpi.com/1424-8220/22/5/1855>.
- 285 [19] I. SPASOJEVIC, V. MURALI, AND S. KARAMAN, *Perception-aware time optimal path parameter-*
286 *ization for quadrotors*, in 2020 IEEE International Conference on Robotics and Automation
287 (ICRA), 2020, pp. 3213–3219, <https://doi.org/10.1109/ICRA40945.2020.9197157>.
- 288 [20] A. TAGLIABUE, D.-K. KIM, M. EVERETT, AND J. P. HOW, *Demonstration-efficient guided policy*
289 *search via imitation of robust tube mpc*, 2022 International Conference on Robotics and
290 Automation (ICRA), (2021), pp. 462–468.
- 291 [21] J. THOMAS, J. WELDE, G. LOIANNO, K. DANILIDIS, AND V. KUMAR, *Autonomous flight*
292 *for detection, localization, and tracking of moving targets with a small quadrotor*, IEEE
293 Robotics and Automation Letters, 2 (2017), pp. 1762–1769, [https://doi.org/10.1109/LRA.](https://doi.org/10.1109/LRA.2017.2702198)
294 [2017.2702198](https://doi.org/10.1109/LRA.2017.2702198).
- 295 [22] J. TORDESILLAS AND J. P. HOW, *MADER: Trajectory planner in multi-agent and dynamic*
296 *environments*, IEEE Transactions on Robotics, (2021).
- 297 [23] J. TORDESILLAS AND J. P. HOW, *PANTHER: Perception-aware trajectory planner in dynamic*
298 *environments*, arXiv preprint arXiv:2103.06372, (2021).
- 299 [24] J. TORDESILLAS AND J. P. HOW, *Deep-panther: Learning-based perception-aware trajectory*
300 *planner in dynamic environments*, IEEE Robotics and Automation Letters, 8 (2023),
301 pp. 1399–1406, <https://doi.org/10.1109/LRA.2023.3235678>.
- 302 [25] C. TOUMIEH, *Decentralized multi-agent planning for multirotors: a fully online and communi-*
303 *cation latency robust approach*, arXiv preprint arXiv:2304.09462, (2023).
- 304 [26] C. TOUMIEH AND A. LAMBERT, *Decentralized Multi-Agent Planning Using Model Predictive*
305

- 306 *Control and Time-Aware Safe Corridors*, IEEE Robotics and Automation Letters, 7
307 (2022), pp. 11110–11117, <https://doi.org/10.1109/LRA.2022.3196777>. Conference Name:
308 IEEE Robotics and Automation Letters.
- 309 [27] R. VAN PARYS AND G. PIPELEERS, *Distributed model predictive formation control with inter-*
310 *vehicle collision avoidance*, in 2017 11th Asian Control Conference (ASCC), IEEE, 2017,
311 pp. 2399–2404.
- 312 [28] Z. WANG, C. XU, AND F. GAO, *Robust trajectory planning for spatial-temporal multi-drone*
313 *coordination in large scenes*, in 2022 IEEE/RSJ International Conference on Intelligent
314 Robots and Systems (IROS), 2022, pp. 12182–12188, [https://doi.org/10.1109/IROS47612.](https://doi.org/10.1109/IROS47612.2022.9982032)
315 2022.9982032.
- 316 [29] X. WU, S. CHEN, K. SREENATH, AND M. W. MUELLER, *Perception-aware receding horizon*
317 *trajectory planning for multicopters with visual-inertial odometry*, IEEE Access, 10 (2022),
318 pp. 87911–87922, <https://doi.org/10.1109/ACCESS.2022.3200342>.
- 319 [30] B. ZHOU, J. PAN, F. GAO, AND S. SHEN, *Raptor: Robust and perception-aware trajectory*
320 *replanning for quadrotor fast flight*, IEEE Transactions on Robotics, 37 (2021), pp. 1992–
321 2009, <https://doi.org/10.1109/TRO.2021.3071527>.
- 322 [31] X. ZHOU, J. ZHU, H. ZHOU, C. XU, AND F. GAO, *EGO-Swarm: A Fully Autonomous and*
323 *Decentralized Quadrotor Swarm System in Cluttered Environments*, Nov. 2020, <https://doi.org/10.48550/arXiv.2011.04183>, <http://arxiv.org/abs/2011.04183> (accessed 2022-07-
324 05). arXiv:2011.04183 [cs] version: 1.
325