

# PMAPREDUCE: A CONFIGURABLE AND INTUITIVE TOOL FOR PARALLELIZING REDUCE IN MAPREDUCE

BRYAN PYO\* AND JUNG SOO CHU†

**Abstract.** MapReduce is a programming model that facilitates concurrent processing and generation of large data using parallel processing of multiple chunks of the data through a map function and a reduce function. In this paper, we introduce the PMapReduce function, which integrates the default parallelized map function from julia in the form of Distributed.pmap, as well as our own implementation of a parallelized reduce, which takes advantage of multiple tiers of reduce layers in order to allow multiple workers to reduce at the same time. We tested our implementation on three different problems: PageRank, Sorting, and TF-IDF. We found that our PMapReduce generally performs faster than the non-parallelized counterparts at the cost of more memory being needed. If the cost of the overhead of using PMapReduce was less than the benefit from the additional parallelization, then there was an overall increase in performance. We found that this was more likely to occur with larger input sizes. Finally, we made an empirical observation that for the sorting problem, the optimal number of parallelized workers stayed fairly consistent regardless of the size of the input sequence to be sorted.

**Key words.** MapReduce, Parallelization, PageRank, Sorting, TF-IDF, Julia

**Code repository:** <https://github.mit.edu/jschu99/18.337-Final-Project>

**1. Introduction.** We are currently in a very digital era, where data is becoming more and more important, as well as our ability to properly process and collect this data. With the abundance of data, parallelization techniques such as MapReduce have also grown in importance. MapReduce is a programming model that facilitates concurrent processing and generation of large data by allowing the parallel processing of multiple chunks of the data. Figure 1 shows the structure of MapReduce. The first of the two main components of a MapReduce problem is a map procedure, which filters and sorts the input data into an easier format for the next step, the reduce method. The next component, the reduce method, then takes these outputs and performs an operation to combine them in a meaningful way into a single output. The benefit of this architecture is that the map operations and the reduce operations can all be done in parallel. Thus, instead of operating on one element at a time, using MapReduce allows multiple elements to be processed at the same time through multiple workers.




FIG. 1. *MapReduce structure*

Julia already has a default, parallelized version of a map function called pmap from the Distributed package. However, the focus of our project is to implement a

\*MIT CSAIL ([bpyo@mit.edu](mailto:bpyo@mit.edu)).

†MIT CSAIL ([jshu99@mit.edu](mailto:jshu99@mit.edu)).

34 parallelized version of the reduce function that uses multiple tiers of reduce operations  
35 in order to perform multiple reduce operations at the same time. We combine the  
36 Distributed.pmap and our parallelized reduce function into a single framework, the  
37 PMapReduce.

38 There are five main classes of MapReduce problems. Metapatterns, organization  
39 patterns, summarization patterns, join patterns, and filtering patterns. We test our  
40 PMapReduce implementation on problems from three of these classes: metapatterns,  
41 organization patterns, and summarization patterns. Metapattern problems aims to  
42 find patterns within patterns within the data. We focus on the PageRank problem,  
43 which tries to measure how important a web page is by counting the number and qual-  
44 ity of links to a page and is introduced in [7]. The central idea of PageRank is that the  
45 more important web pages on the internet are more likely to receive more links from  
46 other websites. For our implementation of PageRank, we use a simple, iterative algo-  
47 rithm that updates weights for each website in the input by analyzing the surrounding  
48 nodes. Organization pattern problems restructure the input data into a more relevant  
49 or easy to use structure. The problem we focus on is a standard sorting problem. The  
50 sorting algorithm we chose divides the array into multiple smaller arrays so that the  
51 reduce workers can individually sort each subarray in parallel. The final reduce step  
52 combines the smaller sorted arrays into the final sorted array. Summarization pattern  
53 problems group similar data to discover new information about the input data, such  
54 as word count. In particular, we focus on the TF-IDF problem, which measures how  
55 important each word in a document corpus is to that corpus, as described in [8]. This  
56 is done by finding the term frequency, which calculates how many times a word ap-  
57 pears in a document and multiplying it with the inverse document frequency, which  
58 measures how many documents a word appears in.

59 We will first go over several related works, especially to our task of integrating  
60 parallelized reduce to Julia. We will then go over details regarding our implementation  
61 and overall design. We will begin by discussing our implementation of PMapReduce  
62 and then discuss the three problems we are testing our implementation on: PageRank,  
63 Sorting, and TF-IDF. We will finally conclude with the results of our experimentation  
64 on the run time and memory allocation of PMapReduce and our overall conclusions.

65 **2. Related Works.** The idea of MapReduce was introduced by Dean and Ghe-  
66 mawat as a way to process large amounts of data, inspired by the map and reduce  
67 functions often used in functional programming [1]. In this paper, they discuss how  
68 the map function can be used to process key-value pairs to generate a set of inter-  
69 mediate set of key-value pairs. A reduce function can then be used to combine this  
70 intermediate set of key-value pairs into a final relevant key-value pair. They discuss  
71 how their implementation can take any problem written in this format and automati-  
72 cally parallelize the problem and execute it across a set of machines. We expand upon  
73 these ideas by allowing inputs that are not necessary dictionaries and using multiple  
74 layers of reduce functions in order to allow for a higher level of parallelization spread  
75 across multiple workers on the same computer instead of across different computers.

76 Since Dean and Ghemawat published their paper on MapReduce, there have been  
77 many papers that explored the use of the MapReduce model to solve a variety of dif-  
78 ferent problems. Li et al. used the MapReduce model in order to process and manage  
79 large-scale datasets in a distributed cluster [6]. They review how it can be used to  
80 generate search indices, perform document clustering, access log analysis, and per-  
81 form a variety of other data analytics. Verma et al. used the MapReduce model in  
82 the field of biology in order to process genetic algorithms [9]. They were able to use

83 Hadoop, an open source implementation of MapReduce, to obtain stable results on  
 84 genetic algorithm problems with up to 100000 variable problems. Finally, Ene et al.  
 85 used the model to process large data in order to perform several different types of  
 86 clustering, specifically k-center and k-median [2]. They were able to discover that  
 87 their MapReduce performance performed equally or better than non-parallel imple-  
 88 mentations and better than other parallel implementations when using a sufficiently  
 89 large dataset.

90 There have also been a few papers that explore implementing MapReduce models  
 91 in Julia. Kavi discusses how they use Julia to implement several fast MapReduce al-  
 92 gorithms to count word frequencies across a large number of documents [3]. Their first  
 93 implementation was done on the CPU using two processes with MPI and their second  
 94 implementation uses a GPU on Julia’s CUDA library. Although, we did not use a  
 95 GPU in our implementation, their implementation of finding the word frequencies was  
 96 helpful in implementing our simpler algorithm for finding word frequencies. Another  
 97 paper by Kourzanov uses a MapReduce model in Julia in order to perform simula-  
 98 tions [4]. In particular, they use it to speed up a Digital Signal Processing (DSP)  
 99 Intellectual Property (IP) model simulation for a Wireless LAN product. They found  
 100 that with 120 workers, the MapReduce model was able to achieve speedups of around  
 101 40x and that with 480 workers, it was able to achieve speedups of around 260x. These  
 102 results were very promising for our own implementation given that they also discuss  
 103 how it was a fairly straightforward implementation of MapReduce for their simulation.

104 There has also been some work on parallelizing calculations in Julia with works  
 105 from people such as Lee et al. that use parallelization to efficiently solve matrix  
 106 calculations [5]. Although certain matrix calculations can be formatted as MapReduce  
 107 problems, they did not use the model for their matrix calculations. We will be taking  
 108 that extra step of transforming our problems of interest into MapReduce problems  
 109 and then further parallelizing the reduce function in Julia for these problems.

110 **3. Design and Implementation.** In this section, we discuss the implementa-  
 111 tion details for PMapReduce and its three application examples, PageRank, Sort, and  
 112 TF-IDF.

113 **3.1. PMapReduce.** Julia already has an implementation for parallelized map  
 114 in the form of Distributed.pmap. It spawns workers that can handle the given map  
 115 task in parallel by splitting the input collection into batches. Its main advantage is  
 116 that it provides an intuitive and easy to use tool for parallel computing, abstracting  
 117 away details such as spawning the workers, distributing tasks, and combining them  
 118 into the single output variable. It offers a multitude of features, such as configurable  
 119 batch size and error handling. The PMapReduce that we implemented is an extension  
 120 of this function, with parallelized reduce integrated as well. The goal is to provide an  
 121 intuitive and easy to use function like Distributed.pmap while integrating parallelized  
 122 reduce.

123 One way to achieve parallelized reduce is by having tiered reduce functions. [Fig-](#)  
 124 [ure 2](#) shows the structure of PMapReduce. Suppose the map function generates  $n$   
 125 items in the output. Traditionally, this can be processed by a single non-parallelized  
 126 reduce function that takes in an input collection and returns a single output. How-  
 127 ever, to parallelize this, we split the  $n$  map output items into  $m_1$  batches ( $n \geq m_1$ )  
 128 and spawn  $m_1$  workers, each of which is running a reduce function on the batch.  
 129 There will be  $m_1$  outputs from this process. The aforementioned reduce workers form  
 130 the first tier of reduce functions. With this design we can add multiple reduce tiers,  
 131 where we let the  $i$ th tier spawn  $m_i$  workers. Then, with  $k$  total reduce tiers, the

132 number of outputs after each reduce tier would be  $m_1, \dots, m_k$ , a non increasing  
 133 sequence. Note that we want the final output to be a single output, so  $m_k$  must be  
 134 1. In terms of performance, this design has benefits and drawbacks. Its main ben-  
 135 efit is that all of the reduce tiers are parallelized, except for the last one. With a  
 136 large input and suitable hardware, this can result in faster performance. However,  
 137 the drawback is that compared to using a single reduce function, the total number of  
 138 operations and memory allocation increase. So, the performance comparison results  
 139 between PMapReduce and the traditional approach of combining Distributed.pmap  
 140 with reduce depend heavily on the scenario.

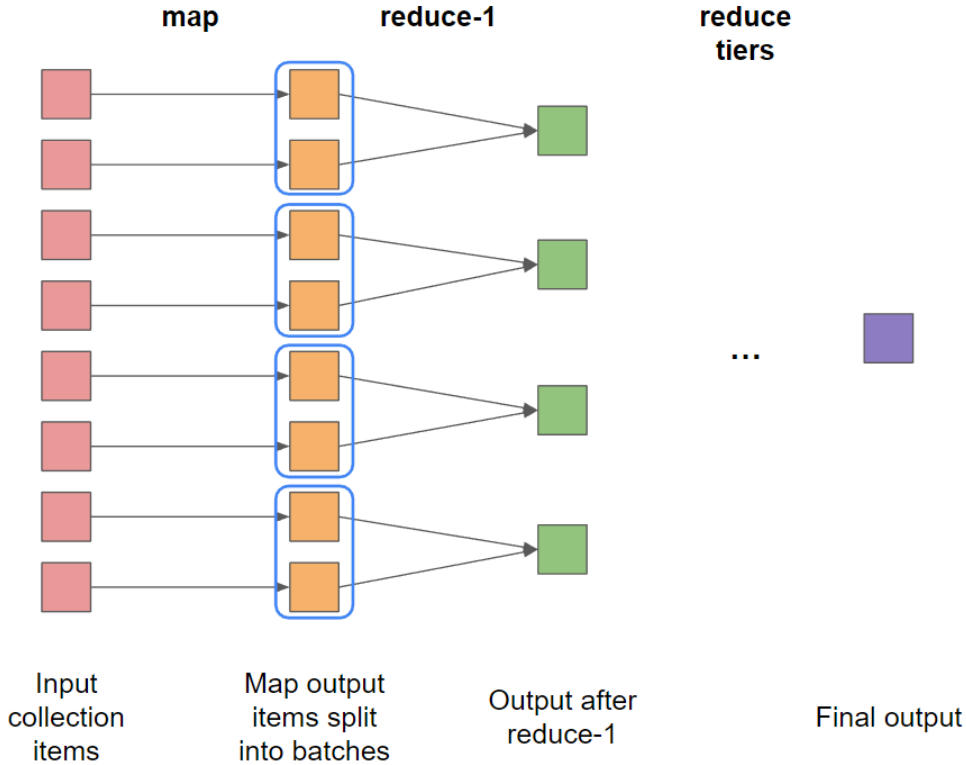


FIG. 2. The structure of PMapReduce with tiered reduce functions

141 Our implementation of PMapReduce takes in five inputs. The function inputs  
 142 were designed so that as much of the implementation needed to run PMapReduce  
 143 would be abstracted away inside the function while still offering a high degree of  
 144 configurability. The five inputs are:

- 145 1. input: This stores the input collection for PMapReduce.
- 146 2. map\_function: This is the map function that will be used by the Distributed.  
 147 pmap part of PMapReduce.
- 148 3. reduce\_functions: This is the collection of reduce functions to be used in the  
 149 parallelized reduce part. Each reduce function is used in a reduce tier.
- 150 4. inter\_results: This is a collection of preallocated collections for storing inter-  
 151 mediate results after each reduce tier. The use of preallocated collections has

152 a performance advantage. This has to be passed in as an input, as opposed  
 153 to automatically being defined inside PMapReduce, because the element type  
 154 depends on the output type of each reduce function.

155 5. `reduce_layer_sizes`: This is a collection of integers for defining the size of the  
 156 outputs after each reduce tier. This is equivalent to the number of workers  
 157 each reduce tier should spawn. Since the last reduce tier should have only  
 158 one worker, so that the final output is size one, `reduce_layer_sizes` should have  
 159 exactly one less element than `reduce_functions` and `inter_results`. For robust-  
 160 ness, if `reduce_layer_sizes` is larger than expected, only the first appropriate  
 161 number of elements are used. Conversely, if it is smaller than expected, the  
 162 missing elements are filled in with ones.

163 **3.2. PageRank with PMapReduce.** PageRank has many versions and dif-  
 164 ferent ways to implement using the MapReduce framework. Since our purpose is to  
 165 compare PMapReduce and the traditional Distributed.pmap and reduce, not to im-  
 166 plement PageRank most efficiently, we use the simple version of it. In this version,  
 167 the weights for each webpage need to converge according to the following formula:

$$168 \quad PR(A) = \frac{1-d}{N} + d \left( \frac{PR(B)}{L(B)} + \frac{PR(C)}{L(C)} + \dots \right)$$

169  $d$  =Damping factor

170  $N$  =Total number of web pages

171  $PR(X)$  =The PageRank of web page  $X$

172  $L(X)$  =The outdegree of web page  $X$

174 Here, damping factor is the probability of the search engine user clicking one of the  
 175 links in the current webpage, as opposed to going to a completely random webpage  
 176 on the internet. We use the standard value of 0.85 for it. According to the above  
 177 expression for PageRank, we can view the weight of a webpage as the sum of weight  
 178 contributions from the webpages that link to it and a constant. Hence, we need to  
 179 compute this efficiently in each iteration. Considering this, the iterative algorithm for  
 computing PageRank is shown in [Algorithm 3.1](#)

---

**Algorithm 3.1** PageRank iterative algorithm

---

Define  $W :=$  weights for the websites

Initialize  $W = \{1/N, \dots, 1/N\}$

**while**  $|W - prev\_W| > t$  **do**

$C := N$  by  $N$  matrix storing weight contributions from web pages to others

$C_{ij} = \mathbb{1}_{ij}dW_i/o_i$ , where  $\mathbb{1}_{ij}$  is 1 if web page  $i$  has a link to web page  $j$ , 0 otherwise.

$d$  is the damping factor,  $W_i$  is the weight of web page  $i$ , and  $o_i$  is the outdegree  
of web page  $i$ .

$new\_W = \{(1-d)/N + \sum_i C_{i1}, \dots, (1-d)/N + \sum_i C_{iN}\}$

$prev\_W = new\_W$

$W = new\_W$

**end while**

**return**  $W$

---

180

181 The iterative algorithm for computing PageRank in [Algorithm 3.1](#) was imple-  
 182 mented using PMapReduce. Specifically, the computation inside the while loop is  
 183 compatible with PMapReduce. [Figure 3](#) shows the implementation in PMapReduce.

184 The computation of  $C_{ij}$  is done using the Distributed.pmap of PMapReduce. Then,  
 185 the reduce tiers compute  $\Sigma_i C_{ij}$ . While many reduce tiers could have been used for  
 186 this, we implemented with just two reduce tiers for the sake of simplicity. The first  
 187 tier computes the sum of a subset of the  $C_{ij}$  rows. For example, if it has  $m$  workers,  
 188 each worker would compute one of  $\Sigma_{1 \leq i \leq N/m} C_{ij}, \dots, \Sigma_{(m-1)N/m \leq i \leq N} C_{ij}$ . Then, the  
 189 second tier adds up the  $m$  outputs from the first tier. With this implementation, it  
 190 is easy to add additional reduce tiers and change the number of workers, which are  
 191 useful when the input size becomes particularly large.

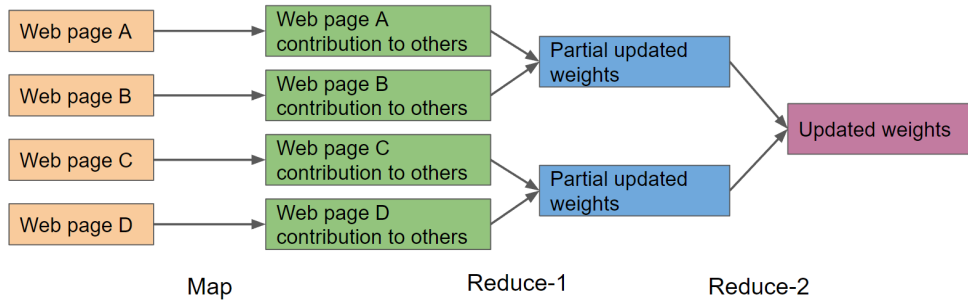


FIG. 3. PageRank implementation using PMapReduce

192 **3.3. Sort with PMapReduce.** While there are many sorting algorithms, only  
 193 some of them can be parallelized without a significant change in the algorithm. We  
 194 chose the sorting algorithm that works with the MapReduce framework, and it is  
 195 shown in Algorithm 3.2. This algorithm can be implemented using PMapReduce  
 196 with two reduce tiers. The Distributed.pmap component splits the input array into  
 197 bins. Then, the first reduce tier spawns multiple workers, each sorting a subset of bins.  
 198 For example, if there are 15 bins,  $A_1, \dots, A_{15}$ , and 3 workers, worker 1 sorts the bins  
 199  $A_1, \dots, A_5$ , worker 2 sorts  $A_6, \dots, A_{10}$ , and worker 3 sorts  $A_{11}, \dots, A_{15}$ . After that,  
 200 the second reduce tier concatenates these sorted bins and produces the final output.  
 201 Figure 4 shows a small example of the sort algorithm implemented with PMapReduce.

---

**Algorithm 3.2** Sorting algorithm

---

Split the input array  $A$  into  $n$  bins:  $\{A_1, \dots, A_n\}$  such that  $\forall i, j \in [1, n]$  such that  
 $i < j$ ,  $\max A_i \leq \min A_j$   
 Sort each bin,  $A_1, \dots, A_n$   
 $A' = \text{concat}(A_1, \dots, A_n)$   
**return**  $A'$

---




FIG. 4. Sort implementation using PMapReduce

202 **3.4. TF-IDF with PMapReduce.** TF-IDF stands for term frequency - inverse  
 203 document frequency, and it is a statistical measure that evaluates how relevant  
 204 a particular word in a document is for a corpus of documents. It is useful in many  
 205 different fields, but especially in automated text analysis and as a method of numerically  
 206 scoring words for natural language processing. It is the result from multiplying  
 207 the term frequency term with the inverse document frequency term.

208 The term frequency measures how often a particular word appears in a particular  
 209 document. Thus, even if the same word appears in the multiple documents, it will  
 210 have its own term frequency value for each document. The term frequency can also  
 211 be calculated in a variety of different ways as long as within the same document,  
 212 the value gets larger with more occurrences. We chose the logarithmically scaled  
 213 frequency method which can be calculated with:

$$TF(term, document) = \log(1 + f_{term,document})$$

214 where  $f_{term,document}$  is simply the raw number of times the *term* appears within the  
 215 *document*. With this logarithmically scaled calculation for term frequency, the value  
 216 increases with more frequent occurrences of the word within the document, but it  
 217 also levels off at some point. This is to show the diminishing returns of importance  
 218 from unique word that appears too many times within a document without being  
 219 particularly useful in measuring relevance, such as certain proper nouns.

220 The inverse document frequency term measures the frequency of a word across  
 221 a set of documents. In other words, it is a way to measure how rare or common a  
 222 word is within a corpus of documents. The closer this term is to 0, the more common  
 223 the word is to the document. This can also be calculated in many ways but the most  
 224 common method way to calculate the term is as follows:

$$IDF(term, document) = \log(1 + \frac{N}{|d \in D : t \in d|})$$

225 where N is the total number of documents within a corpus and  $|d \in D : t \in d|$  is the  
 226 number of documents within a corpus in which the term *t* appears. This term is  
 227 helpful because even if a word appears many times within a single document, it is  
 228 not very relevant within the corpus if it appears in many different documents. For

235 example, words such as “the”, “a”, and “is” are likely to appear many times within a  
 236 single document. However, since these terms also appear across many documents in  
 237 a corpus, it is not very significant relative to a corpus.

238 Our implementation for calculating TF-IDF consists of several layers of mapping  
 239 and reduction. There is first an initial mapping of the corpus to calculate the TF  
 240 values. Once the TF values are obtained, they are then mapped again to represent a  
 241 boolean flag that determines whether or not the word appears in the document. These  
 242 initial boolean values are then reduced through two layers, across multiple workers,  
 243 in order to get a final IDF dictionary for the corpus. Finally, the IDF dictionary as  
 244 well as the previously calculated TF values are combined in order to get the final  
 245 TF-IDF values for all the terms in the corpus. The algorithm is also summarized in  
 246 [Algorithm 3.3](#).

---

**Algorithm 3.3** TF-IDF algorithm

---

Map input document corpus  $\rightarrow$  TF values  
 Map TF values  $\rightarrow$  boolean flags for term existence  
 Parallized reduce boolean flags  $\rightarrow$  IDF dicionaries  
 Final reduction to combine IDF dictionaries  
 $TF - IDF_v\text{alues} = \text{combine IDF dictionary and TF values}$   
**return**  $TF - IDF_v\text{alues}$

---

247 **4. Results.** In order to show the efficacy of PMapReduce, we compared the  
 248 PMapReduce implementations for the three aforementioned problems, PageRank,  
 249 sort, and TF-IDF, with their single worker reduce counterparts. Since we expected  
 250 that the results would heavily depend on the input data size, we experimented with  
 251 varying input size, starting from a very small number until we saw an inflection point,  
 252 where the run time ordering between various implementations changed. Since there  
 253 could be run to run variance, for each datapoint, we ran the experiment 10-20 times  
 254 and used the average. We also collected data on memory allocation, as using memory  
 255 allocation is a major downside for PMapReduce that we anticipated, and it could be  
 256 used to explain its performance behavior.

257 For PageRank, three implementations were compared for their run time and mem-  
 258 ory allocation. The first is an implementation using a single worker map and single  
 259 worker reduce, to provide a baseline implementation. The second is the multiple  
 260 worker map and a single worker reduce. Note that these two do not use PMapReduce,  
 261 as they could be implemented with Julia’s default map, reduce, and Distributed.pmap.  
 262 On the other hand, the last implementation had a multiple worker map and a multi-  
 263 ple worker reduce implemented with PMapReduce. [Figure 5](#) shows the resulting run  
 264 time. Here, the input size represents number of web pages in the input. Each link  
 265 from a web page to the other was put into the input randomly with a 0.3 chance. We  
 266 see that at lower input size, all three implementations had a very similar run time, but  
 267 with larger input sizes, the multi map, multi reduce implementation was the fastest,  
 268 followed by multi map, single reduce and single map, single reduce, in that order. This  
 269 is likely because the three implementations use similar amount of memory allocation,  
 270 as shown in [Figure 6](#), so the parallelization results in a performance benefit.



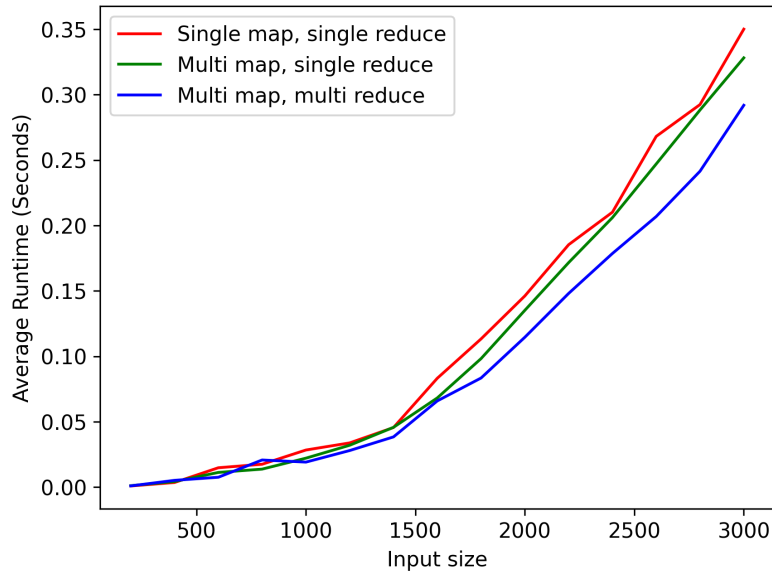


FIG. 5. PageRank run time results




FIG. 6. PageRank memory allocation results

271 The next problem we experimented on was sorting. For this problem, the input  
 272 size represents the number of randomly generated floats in the input array. We used  
 273 them to compare three implementations. The first two were single worker map and  
 274 single worker reduce; and multi worker map and single worker reduce, same as PageR-  
 275 ank implementations. However, for the third implementation that uses PMapReduce,

276 we opted to pair single worker map with the parallelized reduce. This is because, as  
 277 shown in [Figure 8](#), multi worker reduce uses too much memory allocation compared to  
 278 its single worker counterpart. As a result, as shown in [Figure 7](#), the multi map single  
 279 reduce implementation is by far the slowest. So, we have a single map multi reduce  
 280 implementation that we can compare to the single map single reduce implementation.  
 281 According to [Figure 7](#), the two use a very similar amount of memory allocation. So,  
 282 due to the performance benefit with parallelization, single map multi reduce outper-  
 283 forms single map single reduce as expected. In addition to the run time and memory  
 284 allocation analysis, we also experimented with changing the number of workers in the  
 285 reduce tier one of the PMapReduce sort implementation. As shown in [Figure 9](#), we  
 286 experimented with varying input sizes from 10000 to 100000 and number of workers  
 287 from 10 to 150. We observe several interesting trends in [Figure 9](#). First, the number  
 288 of workers can drastically change the average run time. We notice that the runs with  
 289 low number of workers tend to be slow. However, having the highest number of work-  
 290 ers (150) was never the fastest run. Also, the optimal number of workers remained  
 291 fairly consistent at around 125 workers regardless of the input size. While this result  
 292 is hard to generalize, since the data was collected only on a single implementation  
 293 of a problem, it is an interesting empirical result showing the relation between the  
 294 input size and the number of workers, and the optimal number of workers in general.  
 295 Finding the optimal number of workers is a useful step for any implementation using  
 296 PMapReduce.

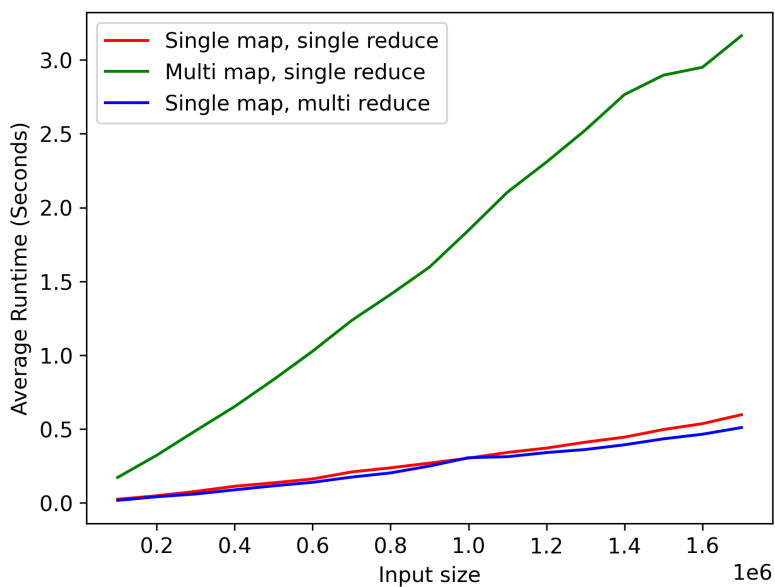


FIG. 7. Sort run time results




FIG. 8. Sort memory allocation results




FIG. 9. Sort optimal number of workers. Red represents faster speed, while black is for slower speed. The fastest run for each input size is highlighted with yellow.

297 The last problem we examined was TF-IDF. Here, the input size represents the  
 298 number of documents. Standard three implementations were used for this problem:  
 299 single worker map and single worker reduce; multi worker map and single worker  
 300 reduce; and multi worker map and multi worker reduce. Only the last implementation

301 used PMapReduce. Also, it is important to note that the TF-IDF implementation had  
 302 three map operations in total, and the multi map single reduce implementation used  
 303 Distributed.pmap for all three, but the multi worker multi reduce implementation  
 304 used Distributed.pmap for only two, and regular map for the other one. This is  
 305 because that one map operation used too much memory allocation when parallelized  
 306 with Distributed.pmap, to a point where it hindered the performance. This is shown  
 307 in Figure 11, where the multi map single reduce uses substantially more memory  
 308 allocation than the other two. As a result, as shown in Figure 10, it had the slowest  
 309 run time. Comparing the single map single reduce and multi map multi reduce, the  
 310 memory allocation of the two are close, although the latter uses slightly more. This  
 311 is as expected, since parallelization and spawning workers involves more operations  
 312 and memory usage. However, the multi map multi reduce ended up being faster  
 313 than the single map single reduce in Figure 10, as the performance benefits from the  
 314 parallelization likely outweighed the slight cost increase in memory allocation.




FIG. 10. *TF-IDF run time results*

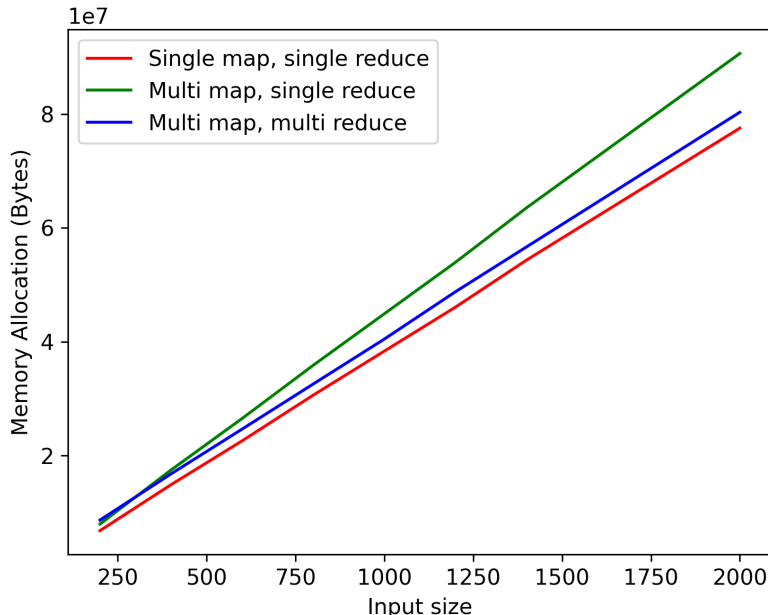


FIG. 11. *TF-IDF memory allocation results*

315 **5. Conclusion.** MapReduce is a common programming model that allows for  
 316 the processing and generation of large data and consists of a mapping function and  
 317 a reduce function. Although Julia already has a parallelized map function in the  
 318 form of Distributed.pmap, in this paper we introduce PMapReduce which can use  
 319 Distributed.pmap as well as our implementation of a parallelized reduce function.  
 320 Our parallelized reduce function uses multiple tiers of reduce functions in order to  
 321 allow multiple workers to work on reducing a set of outputs from the mapping step.  
 322 Every tier except for the last tier of reduce functions can be done in parallel, with  
 323 only the last tier being done sequentially in order to ensure that one correct output  
 324 is generated.

325 Through our experiments, we discovered that in general the parallelized reduce  
 326 method led to faster performance. However, there was also a cost from the over-  
 327 head for implementing our PMapReduce function, which made this less clear. This  
 328 additional cost came from various factors such as creating new workers, deleting work-  
 329 ers, and setting up the correct input and output formats. These cost were observed  
 330 through the increased memory allocation. When the cost of using the PMapReduce  
 331 method was less than the benefit from parallelizing the reduce method, there were  
 332 significant decreases in performance. This disparity was better seen with larger data  
 333 inputs due to parallelization being more impactful than the less scalable cost of setting  
 334 up the problem.

335 We also made an empirical observation that for sorting, regardless of the size of the  
 336 input sequence, the optimal number of threads stayed fairly consistent. While this is  
 337 not a generalizable result, since this data was collected on only one implementation for  
 338 one problem, it was an interesting result. This is likely due to the balance between the  
 339 cost of spawning and assigning tasks to more workers and the benefit from additional  
 340 workers remains stable when the input size is sufficiently large and the hardware does  
 341 not change.

342

## REFERENCES

- 343 [1] J. DEAN AND S. GHEMAWAT, *MapReduce: Simplified Data Processing on Large Clusters*, in  
344 OSDI'04: Sixth Symposium on Operating System Design and Implementation, San Fran-  
345 cisco, CA, 2004, pp. 137–150.
- 346 [2] A. ENE, S. IM, AND B. MOSELEY, *Fast clustering using MapReduce*, Proceedings of the  
347 17th ACM SIGKDD international conference on Knowledge discovery and data mining,  
348 (2011), pp. 681–689, <https://doi.org/10.1145/2020408.2020515>, <https://dl.acm.org/doi/10.1145/2020408.2020515> (accessed 2023-05-16). Conference Name: KDD '11: The 17th  
349 ACM SIGKDD International Conference on Knowledge Discovery and Data Mining ISBN:  
350 9781450308137 Place: San Diego California USA Publisher: ACM.
- 351 [3] N. KAVI, *MapReduce for Counting Word Frequencies with MPI and GPUs*, (2022), <https://doi.org/10.48550/ARXIV.2206.05269>, <https://arxiv.org/abs/2206.05269> (accessed 2023-05-16).  
352 Publisher: arXiv Version Number: 1.
- 353 [4] P. KOURZANOV, *Parallel evaluation of a DSP algorithm using julia*, Proceedings of the 3rd Inter-  
354 national Workshop on Software Engineering for Parallel Systems, (2016), pp. 20–24, <https://doi.org/10.1145/3002125.3002126>, <https://dl.acm.org/doi/10.1145/3002125.3002126> (ac-  
355 cessed 2023-05-16). Conference Name: SPLASH '16: Conference on Systems, Program-  
356 ming, Languages, and Applications: Software for Humanity ISBN: 9781450346412 Place:  
357 Amsterdam Netherlands Publisher: ACM.
- 358 [5] J. H. LEE, Y. KIM, Y. RYU, W. SODSONG, H. JEON, J. PARK, B. BURGSTALLER, AND  
359 B. SCHOLZ, *Julia Cloud Matrix Machine: Dynamic Matrix Language Acceleration on Mul-  
360 ticore Clusters in the Cloud*, Proceedings of the 14th International Workshop on Program-  
361 ming Models and Applications for Multicores and Manycores, (2023), pp. 1–10, <https://doi.org/10.1145/3582514.3582518>, <https://dl.acm.org/doi/10.1145/3582514.3582518> (ac-  
362 cessed 2023-05-16). Conference Name: PMAM'23: 14th International Workshop on Pro-  
363 gramming Models and Applications for Multicores and Manycores ISBN: 9798400701153  
364 Place: Montreal QC Canada Publisher: ACM.
- 365 [6] F. LI, B. C. OOI, M. T. ÖZSU, AND S. WU, *Distributed data management using MapReduce*,  
366 ACM Computing Surveys, 46 (2014), pp. 1–42, <https://doi.org/10.1145/2503009>, <https://dl.acm.org/doi/10.1145/2503009> (accessed 2023-05-16).
- 367 [7] L. PAGE, S. BRIN, R. MOTWANI, AND T. WINOGRAD, *The PageRank Citation Rank-  
368 ing : Bringing Order to the Web*, Nov. 1999, <https://www.semanticscholar.org/paper/The-PageRank-Citation-Ranking-%3A-Bringing-Order-to-Page-Brin/eb82d3035849cd23578096462ba419b53198a556> (accessed 2023-05-16).
- 369 [8] J. E. RAMOS, *Using TF-IDF to Determine Word Relevance in  
370 Document Queries*, 2003, <https://www.semanticscholar.org/paper/Using-TF-IDF-to-Determine-Word-Relevance-in-Queries-Ramos/b3bf6373ff41a115197cb5b30e57830c16130c2c> (accessed 2023-05-16).
- 371 [9] A. VERMA, X. LLORÀ, D. E. GOLDBERG, AND R. H. CAMPBELL, *Scaling Genetic Algorithms  
372 Using MapReduce*, 2009 Ninth International Conference on Intelligent Systems Design and  
373 Applications, (2009), pp. 13–18, <https://doi.org/10.1109/ISDA.2009.181>, [http://ieeexplore.  
374 ieee.org/document/5362925/](http://ieeexplore.ieee.org/document/5362925/) (accessed 2023-05-16). Conference Name: 2009 Ninth Inter-  
375 national Conference on Intelligent Systems Design and Applications ISBN: 9781424447350  
376 Place: Pisa, Italy Publisher: IEEE.