

A High Performance Julia Implementation of Shapely Effects for Global Sensitivity Analysis

Devang Sehgal* and Anurag Vaidya*

Abstract. In this project, we present a high performance Julia implementation of Shapley effects for performing global sensitivity analysis. While many methods exist for sensitivity analysis, they assume independence between the input features of a function, which may be incorrect for many real world scenarios. We are motivated to implement Shapley effects, and make it a part of `GlobalSensitivity.jl`, because they can handle correlated inputs. In this report, we first describe how Shapley effects are calculated using Monte Carlo simulations. Then we show the correctness of our implementation by testing it on commonly used analytical functions and comparing the results with Sobol indices. We show the utility of Shapley effects in handling correlated inputs by considering a Jackson model of a manufacturing plant. We also demonstrate how our implementation can be used on differential equations by considering a dynamic prey-predator system. Finally, we do a thorough performance analysis of the algorithm, and optimize it to be $20\times$ faster than the Python implementations of Shapley effects. Our Julia code is made public at https://github.com/ajv012/shapley_julia.

1. Introduction: Global Sensitivity Analysis. Global Sensitivity Analysis (GSA) is vital for understanding complex systems and model behavior. It measures how the output changes when inputs vary, identifying important inputs and assessing model robustness. This can allow researchers to make more informed use of models in real world scenarios and develop sparser version of models by removing non-essential inputs. Taking derivatives of the output w.r.to the inputs can give a local measure of how much the output changes for a small change in the input. Probabilistic programming provides an alternative approach to sensitivity analysis, by asking how the output of the model changes on average when the input is changed [6]. To motivate the need for GSA, consider a factory receiving multiple orders daily. The output rate of the factory depends on the rate at which six interconnected workstations in the factory work. GSA would tell the factory manager the output rate of the factory is the most sensitive to which workstations, thus helping in managing the operations and logistics. GSA has real world benefits not just in manufacturing, but in studying many systems, like how forest fires spread [7].

`GlobalSensitivity.jl` is the SciML implementation for GSA methods. Some of the commonly used GSA methods which are already implemented in `GlobalSensitivity.jl` include Derivative-based Global Sensitivity Measures (DGSM), Morris method, Sobol's method and Fourier Amplitude Sensitivity Resampling (FAST). Of these, DGSM and Morris method rely on the idea of successively linearizing to approximate the function in question, with Morris taking finite differences to approximate derivatives. Sobol's method, considered a gold standard in GSA, decomposes the variance of model into summands of variances of the input parameters in increasing dimensionality [8] by assuming that the inputs are independent. It is one of the

*Department of Health Sciences and Technology, MIT

40 most widely used methods. `GlobalSensitivity.jl` provides a simple and intuitive interface for
41 conducting GSA on models of interest, which looks like (where f is the model, method is the
42 technique to use):

```
43     res = gsa(f, method, param_range; samples, batch=false)
```

44 Sobol indices are commonly used but assume independence among inputs, which may not
45 hold in real-world scenarios such as physiology-based pharmacokinetic models of the human
46 body's organs. **In contrast, Shapley effects, introduced in cooperative game theory
47 in 1953, do not make any assumptions about input independence and can handle
48 correlated inputs effectively.** The aim of Shapley effects is to determine the contribution
49 of each player fairly in the total payoff achieved by the coalition of players. One notable
50 advantage of Shapley effects is their normalization property, where the sum of effects over
51 all individual variables equals the variance. This normalization property provides a better
52 interpretability in determining the relative importance of variables. Unlike Sobol indices,
53 which cannot be interpreted as percentages of variance, Shapley effects offer this additional
54 insight. Furthermore, while comparing first-order and total order Sobol indices is challenging
55 due to potential differences in their summation values, Shapley effects exhibit a summation
56 property that ensures fairness and avoids assigning excessive importance to a few inputs—a
57 known issue with Sobol indices [7]. The goal of this project is to implement a high performance
58 version of Shapley effects in Julia and add it to `GlobalSensitivity.jl`. The key motivation for is to
59 improve the accuracy of GSA in the case of correlated inputs. Adding computationally efficient
60 version of Shapley effects, which can be used for a wide range of functions, will significantly
61 improve the accessibility of `GlobalSensitivity.jl`. Our high performance implementation of
62 Shapley effects, which is inspired by the R code and pseudo-algorithm provided by the authors
63 of [7] and a Python implementation of Shapley effects [5], is made public at [https://github.
64 com/ajv012/shapley_julia](https://github.com/ajv012/shapley_julia).

65 The report achieves the following:

- 66
67 1. Section 2 describes how Shapley effects are calculated, which involves estimating the
68 incremental cost of adding an input feature of a function to the entire feature set. We
69 describe how Monte Carlo sampling is used to calculate the conditional variances and
70 exemplify the process further by providing code snippets.
- 71 2. The sensitivity of the method to different hyper-parameters is analyzed thoroughly in
72 Section 4.2.
- 73 3. In Section 4.1 the correctness of our implementation of Shapley effects is measured on
74 common analytical functions, like the linear and Ishigami function, routinely used in
75 sensitivity analysis. We also compare the results with Sobol indices and highlight the
76 better interpretability provided by Shapley effects.
- 77 4. We demonstrate the real world utility of Shapley effects by analyzing a Jackson model
78 of a factory system in Section 4.3. We show how Shapley effects can handle correlated
79 inputs better than Sobol indices.
- 80 5. In Section 4.4, we show that in addition to analytical functions, our implementation
81 of Shapley effects can be applied to differential equations. We use the Lotka–Volterra

82 equations to model a prey-predator system and show how Shapley effects can vary
83 over time.

84 6. We perform a thorough performance analysis of our algorithm in Section 5. We show
85 how our implementation of Shapley (optimized on serial and parallel code) is $20\times$
86 faster than a counterpart Python implementation of Shapley effects.
87

88 **2. Implementation of Shapley effects.** Estimating the attribution of features of a func-
89 tion can be thought of as finding the contributions of players in a cooperative game. Tradi-
90 tionally, measures like Sobol indices assume that all players act independently, however this
91 may not be true as their might be interactions between players that need to be accounted for.
92 Shapley effects is a variance based method that assesses the equitable allocation of a player's
93 contribution in a cooperative game by taking into account all of the interactions a player has
94 with others.

95 Before we understand how Shapley effects are calculated, we introduce some notation
96 and definitions. Formally, consider we have a function f which has M inputs represented by
97 $\mathbf{X}_M = \{X_1, X_2, \dots, X_m\}$. An input set \mathbf{X}_J represents a subset of inputs, i.e. $J \subset M$. The
98 marginal distribution of each of the X_i is denoted as D_i , and the joint distribution of any
99 set of input features \mathbf{X}_J is denoted as C_J . The response of the function to the input set is
100 denoted as $Y = f(\mathbf{X}_M)$. The uncertainty in the output Y as a result of \mathbf{X}_M is quantified by
101 $\text{Var}[Y]$ with respect to the joint distribution C_M . Shapley effects quantify how much of the
102 $\text{Var}[Y]$ can be attributed to each of the X_i .

103 The essence of finding the Shapley effect for a feature i is to calculate what is the incre-
104 mental cost of adding the feature i to a subset of features $J \subset M$, which is then averaged
105 over all of the possible sets $J \subset M \setminus i$. In order to systematically consider all the possible
106 subsets of players, we follow these steps:

- 107 1. Consider all possible permutations of the players in the game, denoted as $\Pi(M)$. This
108 will have $m!$ permutations.
 - 109 2. For a permutation $\pi \in \Pi(M)$, define the set $\mathbf{Pi}(\pi)$ as the players that precede player
110 i in Π .
 - 111 3. Define the incremental cost of including player i in $\mathbf{Pi}(\pi)$ as $c(\mathbf{Pi}(\pi) \cup \{i\}) - c(\mathbf{Pi}(\pi))$.
- 112 Now, Shapley effect for a player i (S_i) can be defined as:

$$113 \quad (2.1) \quad S_i = \sum_{\pi \in \Pi(M)} \frac{1}{m!} c(\mathbf{Pi}(\pi) \cup \{i\}) - c(\mathbf{Pi}(\pi))$$

114 In line with previous literature, we define our cost function for any subset $J \subset M$ as:

$$115 \quad (2.2) \quad c(J) = \text{Var}[Y] - \text{E}[\text{Var}[Y|\mathbf{X}_J]]$$

116 The cost function in Equation 2.2 can be understood as the expected reduction in $\text{Var}[Y]$
117 when the values of \mathbf{X}_J are fixed. In other words, how much variance is remaining in Y
118 when the values of \mathbf{X}_J are known. Because of this cost function, we need to evaluate $2^m - 1$

119 variance components for $m!$ permutations of the input features, thus the big-O for the Shapley
 120 algorithm is $m!$, where m is the number of input features. This version of the Shapley effects
 121 algorithm is termed the "exact permutation" version.

122 The exact permutation algorithm quickly becomes intractable for large number of input
 123 features, both in terms of time complexity and the memory allocation requirements. To make
 124 the Shapley effects algorithm tractable for functions requiring a large number of input features,
 125 we implement the "random permutation" version of the algorithm, introduced by Castro et.
 126 al. [1] in 2009. In this algorithm, the essential idea stays the same, but instead of considering
 127 all of the possible $m!$ permutations, we only consider a random subset of all permutations.
 128 It has been shown that such an approximation of the Shapley effects converges to the actual
 129 values in probability [1].

130 **2.1. Pseudo-code for calculating Shapley effects.** To calculate the Shapley effect for
 131 a feature, we need to find the incremental cost of adding the feature to the feature set of
 132 a function, and we define the cost in Equation 2.2. In this section, we describe how our
 133 implementation of Shapley effects can be broken into three main steps. First, we use Monte
 134 Carlo sampling to generate an input sample, X . More specifically, we define distributions
 135 by incrementally adding input features and sample these to encode the interactions between
 136 different features. The input sample X is then passed through the function of interest f ; then
 137 using the definition of cost function and bootstrap sampling, we calculate the Shapley effects
 138 for each feature. We now cover the steps that we follow to generate our input sample X and
 139 provide code snippets for further clarity. To understand how we calculate Shapley effects,
 140 please take a look at our publicly available implementation at [https://github.com/ajv012/](https://github.com/ajv012/shapley_julia)
 141 [shapley_julia](https://github.com/ajv012/shapley_julia).

142 In order to define our input sample X —which encodes the interactions between different
 143 features—we consider all the possible permutations of the features and all of the subsets of
 144 each permutation. Before, we get into the steps we follow to generate X , we outline the
 145 hyper-parameters of our algorithm:

- 147 1. N_V : number of samples used to calculate the variance of the output Y .
- 148 2. N_O : number of samples taken to estimate the conditional variance of the output Y
 149 conditioned on a subset of features \mathbf{X}_J .
- 150 3. N_I : size of each of the N_O samples taken.
- 151 4. n_{perms} : number of permutations of the input features considered. If all permutation
 152 of m features considered, then we have the exact permutation method. In case of
 153 randomly sampling permutations, we have the random permutation method.
- 154 5. n_{boot} : number of bootstrapped samples used to estimate the cost function.

155 To make it easier to follow the steps, consider the following example. Say, we have 5
 156 features, $[1, 2, 3, 4, 5]$, and $m = 5$. The steps we follow to define X are:

- 157
- 158 1. Sample N_V from the joint distribution of all features \mathbf{C}_5 . Store this as sample A.
- 159 2. Consider a permutation of the features $[1, 3, 5, 4, 2]$, call it π .
- 160 3. Within π , select the first feature [1] (call its distribution X^+), and the remaining
 161 features $[3, 5, 4, 2]$ (call their joint distribution X^-).

-
- 162 4. Generate N_O samples of size N_I from the distribution of X^+ conditioned on X^- . Call
163 this set of samples sample B.
164 5. Repeat steps (2) and (3) by incrementally adding the next feature in π to X^+ and
165 keeping the remaining features in X^- . Keep concatenating the generated samples to
166 sample B.
167 6. Repeat steps (2)-(5) for all possible permutations of the features (exact permutation)
168 or a given number of permutation (random permutation).
169 7. Sample A is used to estimate the variance of the output, whereas sample B is used to
170 estimate the mean conditional variance by n_{boot} bootstrapped samples.
171

172 Now we present a code snippet from our implementation that generates sample A and
173 B to encode the correlations between the different features. This is the core of the Shapley
174 algorithm. First, from the method object we extract all of the necessary hyper-parameters.
175 Then, we decide if we are iterating over all permutations of features (exact) or sampling from
176 all permutations (random). We sample N_V from the joint distribution of the marginals. We
177 then have three loops:

- 178
- 179 • First loop iterates over all permutations (step 2 from above),
 - 180 • Second loop samples N_O from X^- ,
 - 181 • Third loop samples N_O samples of size N_I from X^+ conditioned on X^- (step 4).

182 For rest of the algorithm, i.e., how we compute the effects from the input sample, please
183 refer to our public implementation at https://github.com/ajv012/shapley_julia.

```

184
185 if (n_perms==-1)
186     estimation_method = "exact";
187     perms = collect(permutations(range(1,dim), dim));
188     n_perms = length(perms);
189 else
190     estimation_method = "random";
191     perms = [randperm(dim) for i in range(1, n_perms)]
192 end
193
194 # Creation of the design matrix
195 sample_A = copy(transpose(rand(input_distribution, N_V)));
196 sample_B = zeros((n_perms * (dim - 1) * N_O * N_I, dim));
197
198 #---> First loop to go over the permutations
199 for (i_p, perm) in collect(enumerate(perms))
200     idx_perm_sorted = sortperm(perm) # Sort the variable ids
201     for j in 1:(dim-1)
202         # normal set
203         idx_plus = perm[1:j];
204         # Complementary set
205         idx_minus = perm[j+1:end];

```

```

206     sample_complement = sample_subset(input_distribution, N_0, idx_minus);
207     for l in range(1,size(sample_complement)[1])
208         curr_sample = sample_complement[l, :];
209         xj = cond_sampling(input_distribution, N_I, idx_plus, ...
210             idx_minus, curr_sample);
211         xx = reduce(hcat, (xj, repeat(transpose(curr_sample), N_I)));
212         ind_inner = (i_p - 1) * (dim - 1) * N_0 * N_I + (j-1) * ...
213             N_0 * N_I + (l-1) * n_inner;
214         ind_inner += 1;
215         sample_B[ind_inner:ind_inner + N_I - 1, :] = ...
216             @view xx[:, idx_perm_sorted];
217     end
218 end
219 end

```

2.2. Salient features of our implementation of Shapley effects. Our algorithm for calculating Shapley effects has four salient features. First, unlike the R implementation provided by [7], our implementation decouples the sample generation (where different permutations and their subsets are considered to generate conditional samples) and the Shapley effect calculation stages. In the original implementation, since these stages are not decoupled, they cannot be modulated independently. For example, if one wants tighter confidence interval bounds on their Shapley effects, they would need to increase the n_{boot} variable, but this would cause the sample generated to be extremely large as well. This is not necessary as only more samples are needed to make the bounds tighter. This behavior is not seen in our implementation, where tighter bounds on the Shapley effects can be achieved without increasing the computational requirements for the sample generation phase. The second benefit of decoupling these steps is that different high performance techniques can be applied as per the needs of each stage. For example, since the external library <https://github.com/lrnv/Copulas.jl> is heavily used in generating the samples, we cannot parallelize it and can only make improvements to the serial code. In the implementation of [7], this would mean that the Monte Carlo simulations cannot be parallelized. However, with our implementation, we can parallelize the simulations and separately optimize the sample generation code. This is discussed further in section 5. The second salient feature of our implementation of Shapley effects is that it can be applied to both functions with known analytical forms (e.g., linear, Ishigami [2], etc.), as well as a system of differential equations (e.g., Lotka–Volterra equations for a prey-predator system [4]). We analyze both of these systems thoroughly in our experiments in section 3. Third, our implementation allows the users to precisely control the input marginal distributions for all of the features and define a copula to encode the interactions and correlations between the different features. Finally, we implement the exact and random permutation versions of the Shapley effects algorithm so that it is tractable to compute Shapley effects for functions with a larger number of input variables.

3. Experiments. To comprehensively test our implementation of Shapley effects, we perform a series of correctness and exploratory experiments and compare the results with Sobol indices. We also show the utility of Shapley effects. In order to investigate the effect of dif-

249 ferent experimental parameters on estimated Shapley effects—namely N_V , N_O , and N_I —we
 250 use the same systems as in the correctness case.

251 **3.1. Correctness experiments.** To determine the correctness of our implementation of
 252 Shapley effects, we implement commonly used test cases for GSA for which we know what
 253 the expected relative attributions as given by theory. First, we test on the simple linear case
 254 (Equation 3.1), with $A = 7$ and $b = 0.1$, with the assumption that x has a uniform marginal
 255 distribution in the interval $-\pi$ to π . It is expected that feature x would account for 100% of
 256 the variance in the output, so the Shapley effect for x should be close to 1.

$$257 \quad (3.1) \quad y = Ax + b$$

258 Next, we compare the Shapley effects implementation on the Ishigami function [2]. The
 259 Ishigami function (3.2) is a recurrent test case for sensitivity analysis methods and uncertainty.
 260 it is non-linear, non-monotonic, and displays strong inter-dependence between its features,
 261 specifically the first and third as seen in the last term of Equation 3.2. For our tests, we take
 262 all x_i to be uniformly distributed on the interval $-\pi$ to π , and take $a = 7$ and $b = 0.1$. It is
 263 expected that most attribution be given to x_1 and x_2 , and little attribution goes to x_3 because
 264 of its dependence on x_1 .

$$265 \quad (3.2) \quad f(x_1, x_2, x_3) = \sin(x_1) + a \sin^2(x_2) + bx_3^4 \sin(x_1)$$

266 To test the correctness in both linear and Ishigami function, we compare the output of
 267 our Shapley algorithm with that Sobol first and total order indices. We do not expect the
 268 exact attribution values to match due to their different meanings, but expect similar trends in
 269 relative importance given to different features. Additionally, in both the linear and Ishigami
 270 function tests, we also provide the function with an extra feature, which remains unused. The
 271 purpose of this test is to ensure that zero attribution is given to this unused feature.

272 **3.2. Sensitivity to Hyper-parameters.** To determine the effect of hyper-parameters re-
 273 quired for Shapley effects (N_V , N_O , and N_I), we compute the Shapley effects for the Ishigami
 274 function with various combinations of these hyper-parameters and also keep track of the run
 275 time and memory allocations. Trends in these measures will help in determining good trade-
 276 offs between speed, memory, and accuracy of calculated Shapley effects.

277 **3.3. Case study 1: Manufacturing system model.** The first example we consider is
 278 a make-to-order manufacturing system, where we model manufacturing of multiple product
 279 types using a Jackson network [3]. Queueing network models are extensively employed in
 280 industrial engineering and operations research to optimize manufacturing and service systems.
 281 Our objective in conducting sensitivity analysis is to identify the specific product type that
 282 has the greatest impact on fluctuations in the expected order completion time for all jobs.
 283 This valuable information can assist companies in effectively managing system tension and
 284 reducing overall fluctuations. To achieve this, we estimate the Shapley effects and compare
 285 them with the first-order and total effects, providing a comprehensive evaluation of their
 286 respective contributions within this manufacturing system.

287 Consider the network depicted in Figure 1, which represents a manufacturing line with
 288 six workstations labeled A-F, each handling a different job. Throughout a month, the daily

289 arrival rates of six jobs, namely $X_1 \dots X_6$, remain consistent. The model's output is the monthly
 290 expected completion time for the manufacturing line (η), with the \mathbf{X} representing the input
 291 containing the six arrival rates. Since fluctuations in this anticipated time can result in costs
 292 for the company, we use Shapley effects to identify the job types that have the most significant
 293 impact on the variation in expected completion time.




Figure 1. Jackson model of manufacturing line with six workstations (A-F) and six incoming orders (X_{1-6}).

294 Different jobs arrive at different workstations. Job 1 arrives at workstation A, jobs 2, 3, 5
 295 arrive at workstation B, and jobs 4, 6 arrive at workstation E. After each of the job is processed,
 296 they are routed to the next workstation with passing probabilities shown on the arrows in Fig-
 297 ure 1. Job arrival and processing times are assumed to be independent. The processing rates
 298 at each workstation are fixed at: $\mu_A = 1.2, \mu_B = 1.5, \mu_C = 4, \mu_D = 1.8, \mu_E = 3.6, \mu_F = 1.5$
 299 (per day) [7]. Since calculation of Shapley effects requires specification of marginal distribu-
 300 tion of each input, we used Distributions.jl to model each arrival rate with a beta distribution,
 301 $B(\alpha = 1.5, \beta = 2.0, \min = 0.5, \max = 0.8)$.

302 The Jackson model calculates the expected job completion time η given the arrival rates
 303 vector \mathbf{X} , as shown in Equation 3.3. First, the daily arrival rates at each workstation v_{A-F}
 304 are calculated, which are then combined to get the job completion time of the network η .

$$\begin{aligned}
 305 \quad (3.3) \quad & v_A = X_1 \\
 306 & v_B = 0.4X_1 + X_2 + X_3 + X_5 \\
 307 & v_C = 0.3X_1 + 0.15X_4 + 0.15X_6 \\
 308 & v_D = 0.6X_1 + 0.3X_4 + 0.3X_6 \\
 309 & v_E = X_4 + X_6 \\
 310 & v_F = 0.85X_4 + 0.85X_6 + 0.3X_1 \\
 311 \quad & \eta(X_1, \dots, X_6) = \left\{ \sum_{i=A}^F \frac{v_j}{\mu_j - v_j} \right\} \times \left(\frac{24}{\sum_{i=1}^6 X_i} \right) \\
 312
 \end{aligned}$$

313 3.3.1. Numerical experiment: How do correlated inputs affect feature attributions?.

314 We first calculate the Shapley effects of the six order arrival rates assuming that all orders
 315 arrive independently. In line with [7], we use the parameters $N_V = 2,000, N_O = 100, N_I = 2,$
 316 and $n_{boot} = 1,000, n_{perms} = 6!$. In order to make a fair comparison between Sobol indices

317 and Shapley effects, we also implement Sobol first and total order indices so that the same
 318 sampling method used to calculate Shapley effects can be used to find Sobol indices.

319 However, in real world manufacturing and operations research, it is highly unlikely that
 320 all order rates are independent of each other. Let's consider a scenario where products of
 321 types 1 and 2 exhibit a complementary nature, resulting in a positive correlation between
 322 their respective demands. Conversely, products of types 3 and 4 act as substitutes for one
 323 another, leading to a negative correlation in their demands. We now incrementally test the
 324 effect of adding in correlated inputs on the calculated Shapley effects and Sobol indices. First,
 325 we consider "small correlation", where $\text{Corr}(X_1, X_2) = 0.25$ and $\text{Corr}(X_3, X_4) = -0.25$. Next,
 326 we consider "large correlation" where $\text{Corr}(X_1, X_2) = 0.75$ and $\text{Corr}(X_3, X_4) = -0.50$. We
 327 compare the Shapley effects and Sobol indices for both the scenarios.

328 **3.3.2. Numerical experiment: Are all permutations of input features required?.** One of
 329 the limitations of the naive implementation of Shapley effects is that the number of all possible
 330 permutations of the features, n_{perms} , has factorial time complexity. However, all permutations
 331 of features may not be necessary, especially in the scenario with highly correlated inputs. To
 332 make the calculation of Shapley effects tractable for large number of input features, we sample
 333 from all possible permutations. To validate that sampling permutations leads to approximately
 334 similar results, we calculate and compare Shapley effects on the Jackson manufacturing model
 335 with different number of permutations ($6! = 720, 700, 600$, and 200).

336 **3.4. Case study 2: Prey-predatory system of differential equations.** While the first case
 337 study worked with an analytical solution to the Jackson model, in this case study we explore
 338 how Shapley effects can be computed for a system of differential equations. We consider the
 339 Lotka-Volterra equations, which are a pair of first-order nonlinear differential equations and
 340 describe population densities of two inter-connected species, like a prey-predator pair. The
 341 populations change through time according to the pair of equations:

$$342 \quad \frac{dx}{dt} = \alpha x - \beta xy$$

$$343 \quad \frac{dy}{dt} = \delta xy - \gamma y$$

345 Where x and y are the population densities of the prey and the predator, respectively,
 346 $\frac{dx}{dt}$ and $\frac{dy}{dt}$ model the instantaneous growth rates of populations, α is the maximum per
 347 capita growth rate of the prey, and β represents the presence of predators on the growth rate.
 348 Similarly, δ and γ represent the predator's growth rate and presence of prey, respectively.

349 The four parameters are given the following initial values, $\alpha_0 = 1.5, \beta_0 = 1.0, \delta_0 = 3.0, \gamma_0 =$
 350 1.0 . For the sampling of these parameters, we assume uniform distributions over $[1, 5]$. We
 351 solve the problem using the time span of 10 months, and calculate Shapley effects and Sobol
 352 indices at five time points (2, 4, 6, 8, and 10 months) using Tsit5 solver in Julia. Since
 353 the goal of this case study is to show that Shapley effects can be applied to a system of
 354 equations, we assume all parameters are independent. The hyper-parameters for this study
 355 are: $n_{perms} = 4! = 24, N_V = 1000, N_O = 100, N_I = 3, n_{boot} = 1000$.

4. Results.

4.1. Correctness. To test the correctness of our implementation of Shapley effects, we test it on two commonly used functions in sensitivity analysis—linear and Ishigami functions—and compare the attribution given to the input features with Sobol indices. For the linear system (Figure 2A and B), we see that 99.95% of variance is accounted for by the first feature (A in Equation 3.1) and no attribution is given to the second feature (B in Equation 3.1). This is reasonable because the inputs affect the function output via the feature A , and feature B does not account for any of the variance. Similarly, the first order and total order Sobol indices give highest attribution to feature 1. In the Ishigami function, we would expect low attribution for the third feature because it is correlated with the first feature. Shapley effects give 9.8% attribution to the third feature, whereas the remaining attribution is largely given to the first two features (Figure 2C). When interpreting Sobol indices for this example, we see that first order and total order indices demonstrate different stories (Figure 2D). First order indices say that feature 2 is most significant, but total order indices indicate that the first feature is the most significant. Moreover, first order and total order indices cannot be compared against each, so one cannot argue that feature one is most important because total order index assigned to it is the largest out of all Sobol indices (Figure 2D). On the other hand, Shapley effects for different features are comparable because they are percentages of the total variance, hence one can argue that features one and two are more sensitive than feature 3 because they account for a larger percentage of the variation. This highlights the limited interpretation of Sobol indices. Finally, the unused feature, namely the fourth feature, is given the lowest attribution, which is reasonable.

4.2. Sensitivity to hyper-parameters. The Shapley effects algorithm has four hyper-parameters N_V, N_O, N_I, n_{boot} . We analyze the effect of these hyper-parameters on the accuracy, memory allocations, and computation time of Shapley effects, using the Ishigami function as a test case (Table 1). Since N_V controls the size of the sample used to calculate the output variance, it is important not to make that variable too small, even though larger N_V would directly affect the allocation and computation time. However, between 1000 and 10000, we do not see any significant gains but higher increase in costs ($4\times$ increase in memory allocation and $3\times$ increase in computation time). It is also important to balance N_O and N_I because the former denotes the number of samples for calculating the conditional variance and the latter controls the size of the sample. If many small samples are taken ($N_I=2$), we can get non-senseical results like negative Shapley effects. Taking very large samples can cause $7\times$ increase in computation time. Thus, these two variables need to be balanced. Finally, n_{boot} determines the number of bootstrapped samples of size N_V taken; its effect is directly seen on the size of the confidence intervals. We find that 60,000 samples leads to tight confidence intervals, and beyond this, there are more costs than gains. Our selected parameters shown in the last line of Table 1 balance the accuracy, speed, and memory allocation for the Ishigami function. One must note that these values are problem-specific and we encourage users to perform such an analysis to find the best set of hyper-parameters for their problem at hand.

4.3. Case study 1: Manufacturing model. Regardless of the amount of correlation between the features, Shapley effects find that in the Jackson model for the manufacturing line is the most sensitive to feature 2, and this features accounts for large portions of the output

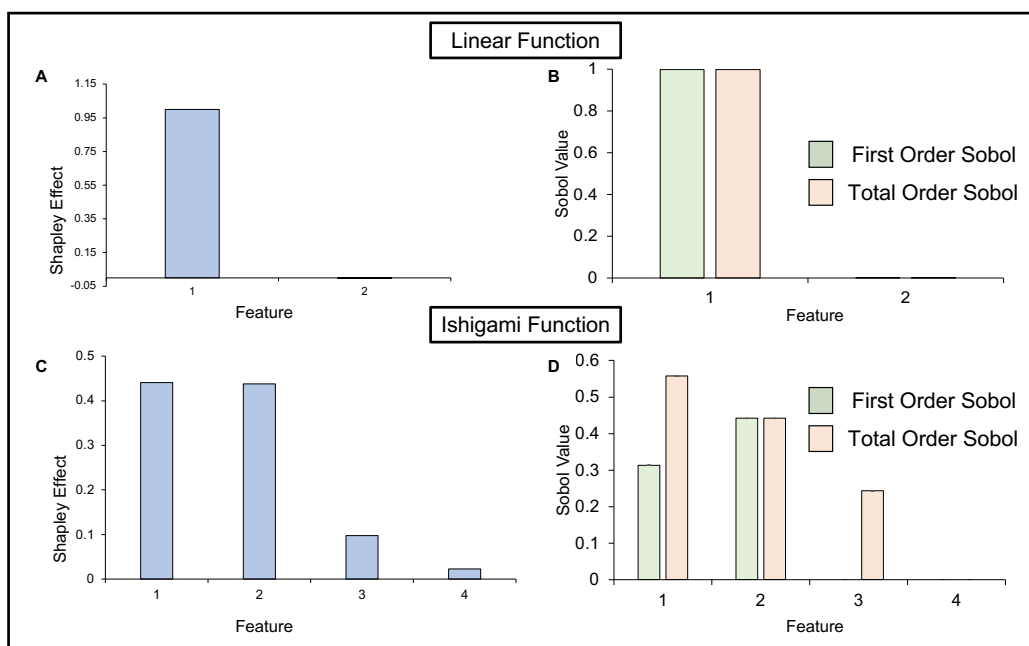


Figure 2. For the linear system described in Equation 3.1 (A) Shapley effects (B) First order and total order Sobol indices. For the Ishigami function described in Equation 3.2 (C) Shapley effects (D) First order and total order Sobol indices. Error bars show 95% confidence interval on n_{boot} samples.

399 variance (Figure 3). We see that as we increase the correlation between the first two fea-
 400 tures, the attributions given to them by the Shapley algorithm also scales accordingly, i.e.,
 401 importance for feature 2 is "routed" to feature 1, which is what we would expect for highly
 402 correlated inputs. On the other hand, for independent and small correlation case, we see
 403 that the Sobol indices also assign highest attributions to the first two features. However, we
 404 once again see that Sobol first order and total order indices are contradicting each other, and
 405 since they are not directly comparable, it is difficult to determine which feature is the most
 406 important. Moreover, we see that as we increase the correlation between the features, Sobol
 407 indices conclude that the first two features are no longer as important as the others, which is
 408 in contrast to the Shapley effects. This helps us validate that when features are correlated,
 409 Shapley effects are a better choice for sensitivity analysis because of fewer discrepancies.

410 Additionally, we use the factory system to determine if all of the permutations of the
 411 features are required to accurately calculate the Shapley effects. This is an important question
 412 because the Shapley algorithm has factorial complexity with the number of permutations,
 413 i.e., the $n_{perms} = (\text{number of features})!$. So, for the factory system, the total number of
 414 permutations to consider are $6! = 720$. For different levels of correlations between inputs,
 415 we find that sampling 80% of total permutations is sufficient to get similar Shapley effects
 416 (Figure 4). Similar attributions are given to the features as low as 200 permutations (28%
 417 of total permutations). Interestingly, we find that the sixth feature gets increased attribution
 418 at lower number of permutations (Figure 4) for both the no correlation and high correlation
 419 cases.

N_V	N_O	N_I	n_{boot}	Feature 1	Feature 2	Feature 3	Allocation (GB)	Time (ms)
10	100	3	60000	0.316 (0.302, 0.329)	0.497 (0.479, 0.515)	0.125 (0.116, 0.135)	1.77	479.86
100	100	3	60000	0.484 (0.483, 0.485)	0.414 (0.413, 0.415)	0.033 (0.033, 0.034)	1.86	493.34
1000	100	3	60000	0.396 (0.395, 0.396)	0.383 (0.382, 0.383)	0.153 (0.153, 0.153)	2.67	631.20
10000	100	3	60000	0.380 (0.379, 0.381)	0.414 (0.413, 0.414)	0.140 (0.139, 0.140)	10.71	1772.0
1000	1	3	60000	0.707 (0.706, 0.708)	0.501 (0.500, 0.501)	-0.253 (-0.254, -0.253)	1.52	333.06
1000	10	3	60000	0.124 (0.123, 0.124)	0.716 (0.714, 0.719)	0.067 (0.065, 0.069)	1.63	351.6
1000	100	3	60000	0.319 (0.318, 0.319)	0.418 (0.417, 0.418)	0.167 (0.167, 0.169)	2.67	579.71
1000	1000	3	60000	0.431 (0.430, 0.432)	0.387 (0.387, 0.387)	0.104 (0.104, 0.104)	12.97	2917.0
1000	100	2	60000	0.459 (0.458, 0.460)	0.402 (0.402, 0.403)	0.075 (0.074, 0.075)	2.66	602.27
1000	100	5	60000	0.382 (0.381, 0.382)	0.394 (0.393, 0.394)	0.157 (0.157, 0.157)	2.67	714.80
1000	100	10	60000	0.459 (0.458, 0.460)	0.395 (0.394, 0.395)	0.083 (0.082, 0.83)	2.67	908.03
1000	100	100	60000	0.396 (0.395, 0.396)	0.448 (0.448, 0.449)	0.096 (0.096, 0.097)	2.69	4684.0
1000	100	3	100	0.388 (0.375, 0.401)	0.413 (0.399, 0.426)	0.123 (0.115, 0.132)	0.03	7.19
1000	100	3	1000	0.437 (0.432, 0.442)	0.412 (0.408, 0.416)	0.078 (0.075, 0.081)	0.07	16.10
1000	100	3	10000	0.436 (0.434, 0.437)	0.379 (0.378, 0.381)	0.107 (0.106, 0.108)	0.48	111.98
1000	100	3	100000	0.392 (0.392, 0.393)	0.440 (0.439, 0.440)	0.079 (0.079, 0.080)	4.42	1029.0
1000	100	3	60000	0.398 (0.397, 0.398)	0.386 (0.385, 0.386)	0.147 (0.146, 0.147)	2.67	605.61

Table 1

The effect of different hyper-parameters N_V, N_O, N_I, n_{boot} on the output of the Shapley effects algorithm on the Ishigami function. The values in bold indicate the hyperparameter that is being changed for that set of experiments. The final row shows the set of hyper-parameters used in making Figure 2C. The Shapley effects for the different features are reported with their 95% Confidence Interval in brackets.

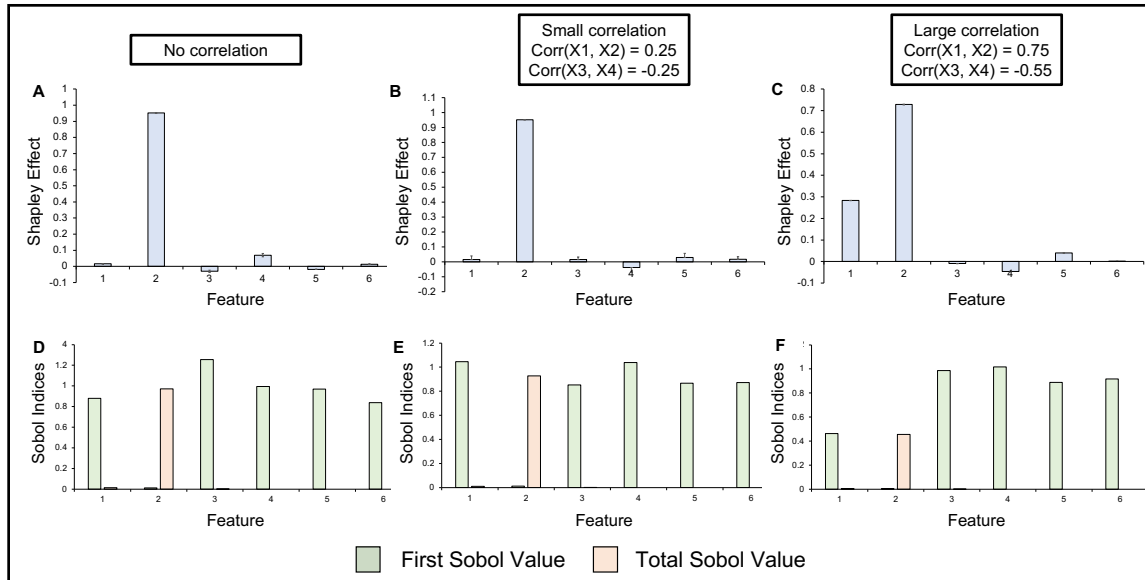


Figure 3. Shapley effects and Sobol indices for the manufacturing line Jackson model explained in Equation 3.3. (A), (B), (C) show Shapley effects for increasing correlations between inputs. (D), (E), (F) show the first order and total order Sobol indices for increasing correlations between inputs. Error bars show 95% confidence interval on n_{boot} samples.

420 **4.4. Case study 2: Prey-predatory system.** The previous case study had an analytical
421 formula that could be analyzed. In this case study, we show that our implementation of
422 Shapley effects can also be applied to a system of differential equations and that the attribution
423 to different features can be analyzed as a function of time. In Figure 5, we see that the different

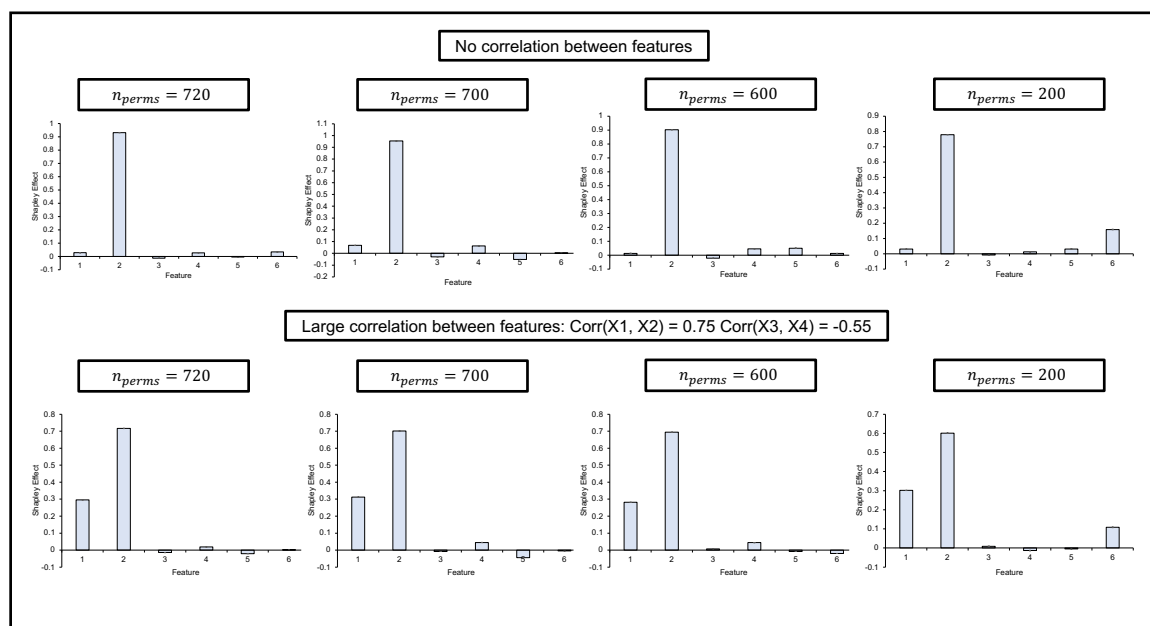


Figure 4. Shapley effects for the manufacturing line Jackson model with different number of randomly selected permutations. We repeat the experiment for when there is no correlation between input features (top row) as well as when there is high correlation between features (bottom row). Error bars show 95% confidence interval on n_{boot} samples.

424 parameters of the prey-predator system have different attributions at different timepoints.
 425 Here, the interpretability of Shapley effects is again exemplified. First, for the prey, we see
 426 that the prey system is most sensitive to the δ parameter (the effect of the presence of prey
 427 on the predator's growth rate), whereas the predator system is generally most sensitive to
 428 the β parameter (the effect of the presence of predators on the prey growth rate). These are
 429 reasonable because the greater growth rate of prey would affect the predatory system (and
 430 vice-versa).

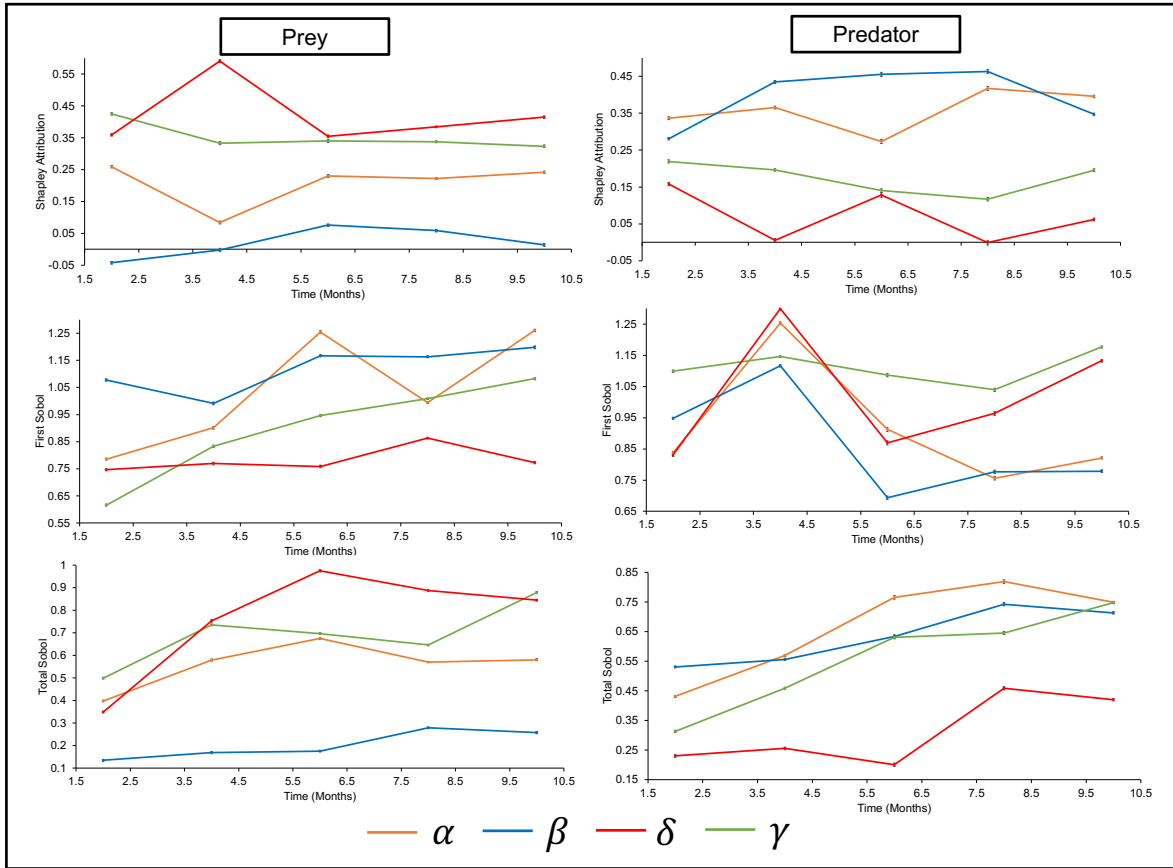


Figure 5. Shapley effects and Sobol indices (first and total order) for the prey-predatory system at different time points.

431 5. Performance Engineering.

432 **5.1. Efforts made in performance improvements.** We started with a serial, naive im-
 433 plementation of the program based on the numpy implementation in [Shapley effects](#) [5]. To
 434 improve performance of our code, we followed the principle of optimising serial performance
 435 first, then looking at parallelism. Our efforts in this can be categorised under three main
 436 steps:

437 **5.1.1. Type Stability.** We first went over our code to ensure all functions were type-stable,
 438 i.e. the compiler knew what type of output and input to expect for every function in the
 439 program. This includes both the core computation GSA function, as well as helper functions
 440 we wrote for sampling. This required going through all the function calls and anticipating what
 441 the type could be, and mentioning that in the function declaration. This did not yield any
 442 noticeable performance improvements over baseline on our test example (Ishigami function).
 443 This is not surprising since our baseline code already had some type-stability incorporated and
 444 the fact that the program has significant computation complexity and requires a large number

445 of allocations which dwarf the effect of type-stability. Nevertheless, our final implementation
 446 of the code is type-stable.

447 **5.1.2. Memory Management (Allocations).** Memory management is extremely impor-
 448 tant as memory requirements can cross 10 GB for even a simple test function like Ishigami
 449 on moderate values of the hyperparamters (see Table 1). Our first effort was in pre-allocating
 450 large arrays before computation to reduce number of allocation calls inside the loops. This
 451 involved use of calls such as

```
452     shapley_indices = zeros(dim, n_boot, 1);
```

453 or `similar` for pre-allocation and then updating elements in loop. This was part of our
 454 baseline implementation.

455 The next step was reducing the number of redundant copying of arrays in loops. This
 456 involved using `@view` where accessing elements of a matrix was required without modification.
 457 In the same vein, it was important to ensure operations were done in-place where possible
 458 with use of Julia’s dot operator (namely `.=`). A very valuable resource for understanding the
 459 major sources of allocations and finding whether our if our operators were indeed allocating
 460 memory as we expected was the use of Julia’s track allocations command line feature, namely
 461 `--track-allocation=user` flag. This generated a `.mem` file listing the number of allocations
 462 per line of code, which was a big help in identifying the key bottlenecks. This flag, along
 463 with BenchmarkTools such as `btime` helped us see that most allocations were occurring inside
 464 the bootstrap loop (apart from our expected array allocations). Taking a closer look at the
 465 bootstrap loop, two lines are of particular importance:

```
466     idx_for_var = rand(1:n_var, n_var);
```

```
467     idx_for_cond_var = rand(1:n_outer, n_outer);
```

468 These lines are used for ‘creating’ a sample for cost computation by selecting elements
 469 from our generated samples. They involve creating an N_V and N_O element random vector at
 470 each iteration of this loop, i.e. resulting in Order ($n_{boot} * N_V + n_{boot} * N_O$) allocations, a very
 471 significant cost considering $n_{boot} = 60,000$ and $N_V = 1000$ are typical and have been used
 472 in this report. The allocations can be brought down by making these allocations in-place,
 473 i.e. allocating the array once and storing new random elements to it each loop instead of
 474 allocating more space. This would look as follows:

```
475     rand!(idx_for_var, 1:n_var );
```

```
476     rand!(idx_for_cond_var, 1:n_outer);
```

477 However, we could not simply make all function calls in-place as we wanted to parallelize
 478 this loop across multiple threads. There was a trade-off to consider between reducing alloca-
 479 tions and reducing computation time by multi-threading. To better understand its effect on
 480 performance, we developed both an optimised serial code that reduced allocations as much as
 481 possible (by using `rand!()` above), and an optimised parallel code designed to make use of
 482 the threading but at the cost of increased allocations.

483 **5.1.3. Multi Threading/ Parallelism.** Use of reductions: Our implementation relies on
 484 generating samples, calculation of the cost function based on the samples and finally calcu-
 485 lating the Shapley values. All these steps require construction, storage, and manipulation of
 486 large arrays. For efficient construction and concatenation we make use of the `reduce` function
 487 in Julia with `hcat` and `repeat` to manipulate our arrays in desired shape. An example is seen

488 as follows:

```
489     xx = reduce(hcat, (xj, repeat(transpose(curr_sample), n_inner)));
```

490 Multi threading: Threading is the most obvious use of parallelism in the code, and we
 491 also expect it provide the most significant benefit in performance. The most straightforward
 492 application is using the `Threads.@threads` call on for loops. We tackle this in two parts,
 493 representing the two key segments of our code:

494 A. Sample generation: Sample generation involves listing the permutations of the dimen-
 495 sions, then generating samples based on conditional distributions and storing them in the
 496 input sample array. The sample generation segment has multiple nested loops with function
 497 calls to our sampling/distribution helper functions (one of which itself contains loops), so it
 498 represents an opportunity for speedup. However, use of threads do not actually improve the
 499 performance of our code - at best it simply reproduces the performance of serial code, or
 500 only slightly worsens it, when the outermost loop is threaded. Threading the innermost loop
 501 significantly worsens the code performance by as much as a factor of 2.

502 We believe this is down to two main reasons. Firstly, we know generating threads has
 503 a significant overhead involved, and memory sharing across threads can lead to slowdowns.
 504 This is especially true when the same data structure is being accessed by multiple threads at
 505 once, which is true in this case - `sample_B` is being used by all threads for storing results while
 506 `input_distribution`, a `sklar` distribution type, is being passed to functions by all threads.
 507 Secondly, the function calls and data types used in this section includes an external library
 508 called Copulas. Since we do not have complete control over all aspects in these loops, we are
 509 unable to control it's behaviour and modify it to be appropriate for multi-threading (discussed
 510 previously in Section 2.2). Therefore, we keep the sample generation part of the code as a
 511 serial implementation but optimise it as well as possible in terms of memory allocations.

512 B. Shapley indices calculations: This segment of the code has two sets of loops. The first
 513 is a single for loop over the n_{boot} for computing the cost function, while the second is two
 514 nested for loops over n_{boot} and number of permutations that computes and stores the actual
 515 Shapley indices. For this part, we use threading across both the outer loops. This gives us a
 516 significant benefit and helps optimize the performance of the code. However this benefit must
 517 be traded-off against a cost in allocations, discussed below in performance analysis.

518 Other libraries and tools beyond `Threads.@threads` were tried as well (such as `LoopVec-`
 519 `torisation.jl`) but no noticeable benefit over the `Threads.@threads` was seen. An additional
 520 benefit of using `Threads` in this manner was its versatility whereas several other tools had
 521 more specific requirements from underlying code.

522 **5.2. Benchmarks and Performance Analysis.** For analysing the performance of our pro-
 523 gram, we measure it's performance using the 3 dimensional Ishigami function with $N_V = 1000$,
 524 $N_O = 100$, $N_I = 3$, $n_{boot} = 60,000$. We benchmark: (i) the python Shapley effects implemen-
 525 tation [5], which served as the inspiration for our code, (ii) our baseline Julia implementation
 526 in the structure of the GSA repo, (iii) our optimised serial implementation for the Julia code,
 527 and (iv) the optimised parallel implementation for the Julia code. The runtime and mem-
 528 ory allocations for the implementations is shown in Table 2. Note that the allocations for
 529 python code are not listed as no direct analogue to `btime` or Julia's `BenchmarkTools` could be
 530 found and the python functions for similar functionality may have differences. All benchmarks

531 reported are for a 2020 M1 processor Macbook Air, with 4 threads in case of multi-threading.

Implementation	N_V	N_O	N_I	n_{boot}	dim	Runtime (s)	Allocations (GB)
python-shapley [5]	1000	100	3	60000	3	12.009	n/a
julia-baseline	1000	100	3	60000	3	1.917	4.12
julia-optim. serial	1000	100	3	60000	3	1.372	1.99
julia-optim. parallel	1000	100	3	60000	3	0.582	2.67

Table 2

Performance benchmarks, namely runtime and allocations, for 4 implementations of Shapley effects for GSA: 'python-shapley' refers to the implementation of Shapley effects in python [5], 'julia-baseline' is our implementation of Shapley-effects in the GSA scaffold, 'julia-optim. serial' refers to our optimised serial implementation while 'julia-optim. parallel' is our optimised parallel implementation, running over 4 threads.

532 We see that the baseline Julia implementation is itself $6.3\times$ faster than the Shapley effects
 533 implementation in python, underscoring the benefits of developing high-performance code in
 534 Julia. Our optimised serial implementation is about $1.4\times$ faster than the baseline, while
 535 using less than half the memory, 1.99 GB vs 4.12 GB, representing a significant saving. Our
 536 optimised parallel code comes in at $3.3\times$ faster than the baseline and $2.4\times$ faster than the
 537 optimised serial code. However, its performance comes at a cost - the parallel code uses 35%
 538 more memory than the optimised serial code. Overall, our fastest parallel code is over 20x
 539 faster than the existing equivalent python implementation, representing a significant markup
 540 in speed.

541 As mentioned, the serial code is significantly more memory efficient, 1.99 GB vs 2.67 GB
 542 for the parallel implementation. This is down to two factors - an increase in allocations and
 543 overhead due to multi-threading, and the fact that we can use in-place operations in the
 544 bootstrap loop. Depending on the problem and system used, either of time or memory may
 545 be more important to optimise over. In this case, since a sizeable time benefit is seen on just 4
 546 threads, parallelization may be preferred. This is likely to be the case for most systems since
 547 HPC setups typically have many more threads. Moreover, while the allocations are increased,
 548 these are in temporary variables so memory is flushed at end of loop; the risk of going out
 549 of bounds in memory are limited. However, there may be certain cases where memory is
 550 more critical and for this reason we publish both our optimised serial and parallel code on the
 551 project's GitHub.

552 For analysing performance of the program, we make use of Table 1, computed using our
 553 optimised parallel code on the Ishigami system for different values of N_V, N_O, N_I, n_{boot} . In
 554 addition, we benchmark our program for different values of dimensions and n_{boot} on Ishigami,
 555 also recording the breakup of the time spent on sample generation vs Shapley indices calcu-
 556 lation, as shown in Table 3.

557 From Table 1, we see the most crucial variable influencing performance of the program
 558 is n_{boot} . Computation time and memory scale roughly linearly (though not exactly, with
 559 variation at both highest and lowest values) with n_{boot} . This follows our expectation since
 560 most of the computation effort and allocations are in the loop going over n_{boot} . n_{boot} sets
 561 the size of the confidence interval, therefore the user must make this key tradeoff between
 562 computation cost and precision. This, and other considerations relating to setting the values

563 of N_V, N_O, N_I were discussed in Section 4.2.

564 Table 3 shows that sample generation time is independent of bootstrap runs. It depends
 565 primarily on the dimension of the problem since that defines the number of permutations over
 566 which the samples are to be generated. This is a key feature of our implementation (discussed
 567 in Section 2.2), distinct from the Song Nelson Staum paper [7] that inspired this work. Table
 568 3 also shows how increasing dimension drastically increases memory and time requirements
 569 in both sections of the code. This underscores the importance of the random permutation
 570 implementation - generating all samples using the exact permutation method for even 10+
 571 dimensions may prove to be very costly. Fundamentally, the use of random permutations
 572 method makes the problem of Shapley effects calculation tractable.

N_V	N_O	N_I	n_{boot}	dim	Total runtime (s)	Total allocations (GB)	Sample time (s)	Shapley time(s)
1000	100	3	1000	3	0.014	0.073	0.006	0.008
1000	100	3	1000	4	0.075	0.281	0.038	0.038
1000	100	3	1000	5	0.531	2.02	0.370	0.161
1000	100	3	1000	6	3.590	13.71	2.162	1.428
1000	100	3	1000	7	45.083	128.35	22.903	22.180
1000	100	3	100000	3	1.054	4.42	0.006	1.048
1000	100	3	100000	4	3.387	15.45	0.081	3.305
1000	100	3	100000	5	20.629	88.19	0.292	20.337

Table 3

Runtime and allocations as a function of number of dimensions and number of bootstraps. The total runtime is broken down into the time taken for generating the sample (sample time) and the time for calculating Shapley indices (shapley time).

573 **6. Conclusion.** In this report, we discussed the importance of global sensitivity analysis,
 574 a crucial technique for analysing the behaviour of a system over entire range of its input pa-
 575 rameters. We discussed the motivation for developing an implementation for Shapley effects,
 576 a GSA technique which is interpretable and can account for dependencies in input parameters
 577 of a system. We described our implementation of Shapley effects in Julia. This involved esti-
 578 mating the incremental cost of adding an input feature of a function to the entire feature set,
 579 and using Monte Carlo sampling to calculate the conditional variances. Our implementation
 580 decouples the generation of samples from the computation of Shapley indices, allowing for
 581 independent modulation of these two segments. It follows the SciML principles and can be
 582 applied to functions and differential equations alike. Another key feature is that it allows
 583 using a random subset of permutations for generating samples, making estimating Shapley
 584 effects tractable in higher dimensional systems as well. We demonstrated the correctness of
 585 our system by testing its result on the Ishigami and Linear functions and comparing with
 586 Sobol. We demonstrated the benefit of our implementation of the Jackson model of a fac-
 587 tory system, showing how the sobol estimates can differ widely from Shapley in the case
 588 of dependent inputs. We also applied our program to a prey-predator system modelled by
 589 Lotka-Volterra ODEs to underscore its versatility. Finally, we optimised our code for per-
 590 formance by focusing on type-stability, reducing allocations and parallelism. We showed our
 591 optimised parallel code running on 4 threads is $20\times$ faster than the equivalent Shapley effects
 592 implementation in python, while our optimised serial code in julia is about $9x$ faster than the
 593 python implementation. Our optimised code is available on the project [repository](#).

594 **Acknowledgments.** We would like to thank Vaibhav Dixit from the Julia hub team for
595 providing valuable feedback in implementation and our approach.

596

REFERENCES

- 597 [1] J. CASTRO, D. GÓMEZ, AND J. TEJADA, *Polynomial calculation of the shapley value based on sampling*,
598 *Computers & Operations Research*, 36 (2009), pp. 1726–1730.
- 599 [2] T. ISHIGAMI AND T. HOMMA, *An importance quantification technique in uncertainty analysis for computer*
600 *models*, in [1990] Proceedings. First international symposium on uncertainty modeling and analysis,
601 IEEE, 1990, pp. 398–403.
- 602 [3] J. R. JACKSON, *Jobshop-like queueing systems/comments on" jobshop-like queueing systems"*, *Management*
603 *science*, 50 (2004), p. 1796.
- 604 [4] Y. LI AND Y. KUANG, *Periodic solutions of periodic delay lotka–volterra equations and systems*, *Journal*
605 *of Mathematical Analysis and Applications*, 255 (2001), pp. 260–280.
- 606 [5] B. NAZIH, *Cemracs17 / shapley-effects*, 2017, [https://gitlab.com/CEMRACS17/shapley-effects/-/tree/
607 master/](https://gitlab.com/CEMRACS17/shapley-effects/-/tree/master/).
- 608 [6] C. RACKAUCKAS, *Sciml/scimlbook: v1.1*, Nov. 2022, <https://doi.org/10.5281/zenodo.7347643>.
- 609 [7] E. SONG, B. L. NELSON, AND J. STAUM, *Shapley effects for global sensitivity analysis: Theory and*
610 *computation*, *SIAM/ASA Journal on Uncertainty Quantification*, 4 (2016), pp. 1060–1083, [https:
611 //doi.org/10.1137/15M1048070](https://doi.org/10.1137/15M1048070).
- 612 [8] X. ZHANG, M. TRAME, L. LESKO, AND S. SCHMIDT, *Sobol sensitivity analysis: A tool to guide the devel-*
613 *opment and evaluation of systems pharmacology models*, *CPT: Pharmacometrics & Systems Pharma-*
614 *cology*, 4 (2015), p. 69–79, <https://doi.org/10.1002/psp4.6>.