

1 **18.337 FINAL PROJECT**  
2 **GPU RAY TRACING AND PLANETARY RADIATIVE TRANSFER**

3 ZEYAD AL AWWAD \*

4 **Abstract.** This report focuses on high-performance GPU ray tracing, implemented entirely  
5 within Julia using a custom CUDA kernel. It discusses two implementations targeting very dif-  
6 ferent use cases: the first focuses on rendering 3D images, while the second focuses on simulating  
7 radiative transfer (heating/cooling) between an Earth-like planet and Sun-like star. The results  
8 include rendered images, plots of ice-albedo feedbacks in radiative transfer, and performance analy-  
9 sis for all of the main subcomponents (examining how the runtime scales with the number of rays  
10 and objects). A preliminary ray tracing code is available at [https://github.com/Zeyad-Awwad/Julia-](https://github.com/Zeyad-Awwad/Julia-Ray-Tracing/tree/main)  
11 [Ray-Tracing/tree/main](https://github.com/Zeyad-Awwad/Julia-Ray-Tracing/tree/main), though I hope to develop this concept much further over the coming months  
12 (and since I'm just starting to prepare my thesis proposal, most likely years).

13 **1. Introduction.**

14 Coming from a computational physics background, and working for several years  
15 with satellite data and optical astronomy, I am particularly interested in using high-  
16 performance GPU ray tracing to study exoplanet compositions. Almost everything we  
17 can observe about other planets is optical: we can see how much light they receive,  
18 how much they absorb, how much they reflect, how atmospheres interact with the  
19 stellar spectrum, and so much more. The current leading methods for exoplanet  
20 analysis are transit methods (which focus on changes in brightness over a 3D orbit)  
21 and The complex 3D structure of exoplanets is an active area of research, but most of  
22 the spectral analysis of exoplanet observations still relies on simplified 1D models that  
23 ignore several important but challenging components (especially clouds) that might  
24 be addressed through ray tracing.

25 The paper is organized with the following high-level structure: the methodologi-  
26 cal details of my implementations are described in [section 2](#), some preliminary results  
27 are shown in [section 3](#), performance testing (with a focus on scaling) is in [subsec-](#)  
28 [tion 3.3](#), and some of my plans for future work that I hope to include within my thesis  
29 [subsection 4.1](#).

30 **2. Methodology.**

31 **2.1. Ray Tracing.**

32 Although I developed two different models for this project (one for rendering,  
33 one for radiative transfer between a planet and star), both rely on the same ba-  
34 sic formulation. The first section focus on general ray tracing as a computational  
35 workhorse, rather than the details of any specific implementation. It focuses on how  
36 rays are parameterized, how intersections are solved with implicit geometries, and  
37 how ray bounces are handled. The renderer uses all of these features, while the ra-  
38 diative transfer model only uses a subset (and adds several features that are very  
39 physics-specific).

40 The subsequent sections describe some features that are specific to rendering (such  
41 as color tracking and pixel averaging) and to atmospheric physics, which (even in a  
42 simple form) involve plenty of additional physical properties that must be defined,  
43 tracked and updated.

---

\*PhD Candidate in Computational Science and Engineering  
Engineering Systems Laboratory, Department of Aeronautics and Astronautics  
MIT

### 2.1.1. Parameterization and Intersections.

**Primary references:** [3] for spherical geometries, [2] for triangular geometries.

Each ray is represented by a parameterized equation, defined by an origin  $\vec{x}_0$  and a direction  $\vec{d}$

$$\vec{x}(t) = \vec{x}_0 + t \cdot \vec{d}$$

This form allows us to define where a ray intersects a surface by solving for the parameter  $t$  by representing geometries with the appropriate implicit equations. The first positive intersection determines where (and from which object) the ray will bounce, and can be defined mathematically as  $t_i = \min\{t_{i,j} \mid t_{i,j} > 0\}$  where  $t_{i,j}$  defines the intersection point between ray  $i$  and object  $j$ .

The standard approach to ray tracing uses implicit geometries, where the intersection occurs at some solution  $t$  to that equation (or set of equations). Most software uses triangular meshes to represent arbitrary geometries, but solutions to perfect spheres can also be calculated very efficiently using this approach. In practice, these two choices represent the main two geometries considered in real-world ray tracing code, but in principle, other types of geometries can be represented in similar ways.

Spheres have a relatively simple solution, since you essentially just have to solve a quadratic equation where the coefficients are given by a set of dot products between the ray and the sphere (specifically its radius  $\vec{r}$  and centroid  $\vec{c}$ )

$$t^2(\vec{d} \cdot \vec{d}) + t(2\vec{c} \cdot \vec{d}) + \vec{c} \cdot \vec{c} - \vec{r} \cdot \vec{r}$$

This can be solved using the standard quadratic formula. Both solutions need to be checked for validity, since only real  $t$  correspond to actual intersections (complex  $t$  indicates no intersection at all) and only positive  $t$  corresponds to intersections in the direction of the ray (negative  $t$  indicates that the ray would intersect if we reverse the direction).

$$t = \frac{-(2\vec{c} \cdot \vec{d}) \pm \sqrt{(2\vec{c} \cdot \vec{d})^2 - 4(\vec{d} \cdot \vec{d})(\vec{c} \cdot \vec{c} - \vec{r} \cdot \vec{r})}}{2(\vec{d} \cdot \vec{d})}$$

Ray-triangle intersections are more computationally intensive, especially working in Cartesian coordinates (where each ray would have to solve a  $3 \times 3$  system of equations). The Moller-Trumbore method is a popular method that uses a barycentric coordinate transformation to compute these intersections more efficiently.

For a triangle with vertex coordinates  $\vec{A}$ ,  $\vec{B}$  and  $\vec{C}$ , the transformation involves computing 3 translations with respect to vertex  $\vec{A}$

$$\vec{T} = \vec{x}_0 - \vec{A}$$

$$\vec{E}_1 = \vec{B} - \vec{A}$$

$$\vec{E}_2 = \vec{C} - \vec{A}$$

And two cross-products that transform an arbitrary triangle into an axis-aligned unit triangle

$$\vec{P} = \vec{d} \times \vec{E}_2$$

$$\vec{Q} = \vec{T} \times \vec{E}_1$$

82 This allows us to compute  $t$  with a more straightforward and GPU-friendly vector  
83 calculation

$$\begin{bmatrix} t \\ u \\ v \end{bmatrix} = \frac{1}{\vec{P} \cdot \vec{E}_1} \begin{bmatrix} \vec{Q} \cdot \vec{E}_2 \\ \vec{P} \cdot \vec{T} \\ \vec{Q} \cdot \vec{D} \end{bmatrix}$$

84 Since  $t$  only determines when the ray intersects the plane of the triangle, we need  
85 to use the values of  $u$  and  $v$  to determine whether or not the ray is contained within  
86 the boundaries of the triangle. In Cartesian coordinates this would require a more  
87 involved point-in-polygon check, but in barycentric coordinates (where the triangle is  
88 axis-aligned with unit lengths, except for the hypotenuse) it's sufficient to check that  
89  $0 \leq u \leq 1$ ,  $0 \leq v \leq 1$  and  $0 \leq u + v \leq 1$ .

90 Refraction is primarily handled using Snell's law:  $\sin(\theta_2) = \frac{\eta_1}{\eta_2} \cdot \sin(\theta_1)$ , for two  
91 materials with refractive indices  $\eta_1$  and  $\eta_2$ . In cases where  $\frac{\eta_1}{\eta_2} \cdot \sin(\theta_1) > 1$ , this doesn't  
92 have a solution and we instead revert to specular reflection (physically, this represents  
93 the case of looking at water or glass surfaces from a nearly parallel angle).

#### 94 2.1.2. Bounces.

95 Once we find the location of the first positive intersection, we update the position  
96 of the ray and decide on a new direction. The method used to determine the new  
97 direction is material-dependent and represent different physical phenomena such as  
98 specular reflections, diffuse light, refraction and thermal absorption.

99 Specular reflections represent an idealized "perfect mirror", where the angle of  
100 incidence (relative to the surface normal  $\vec{n}$ , of unit length) matches the angle of  
101 reflection. The updated direction is calculated using

$$\vec{d}_{new} = \vec{d} - 2(\vec{n} \cdot \vec{d})\vec{n}$$

102 Diffuse reflections draw a random new direction from some distribution. The  
103 simplest approach is to use a spherically symmetric distribution, but the Lambertian  
104 distribution is more physically accurate (since light is more likely to reflect in the  
105 perpendicular direction. Samples are drawn uniformly from the surface of a unit  
106 sphere that makes contact with the surface at the intersection point (using rejection  
107 sampling from a  $2 \times 2 \times 2$  cube to find points inside the sphere, then normalizing to  
108 unit length to get points on the surface). The vector between the intersection point  
109 and the sampled surface defines the direction of the bounced ray.

#### 110 2.2. Rendering.

111 To capture images using the traditional ray tracing code, we need to cast the  
112 rays from the camera rather than a light source. Each pixel emits the same number  
113 of rays directly perpendicular to the camera plane, with randomly sampled initial  
114 positions within the pixel boundaries. By tracking the colors of the rays, and taking  
115 the average, we can determine the color seen by each pixel.

116 The color of each ray is represented by an RGB vector, initially at  $[1.0, 1.0, 1.0]$   
117 for all rays. With each bounce, the ray loses some color based on the material. A  
118 grassy surface, for example, may absorb 50% of the red and blue channels but leave  
119 the green channel untouched, so the bounced ray would have an updated RGB vector  
120 of  $[0.5, 1.0, 0.5]$  (and a second bounce would give  $[0.25, 1.0, 0.25]$ , and so on).

121 A ray terminates when it flies off into the sky, the primary light source (which has  
122 a uniform blue value, and casts a slight blue tint on all visible objects). When this  
123 occurs, or when all the color values diminish too low to matter very much, their color

124 is added to the original pixel they were cast from. This is divided by the number of  
 125 rays to get an average color value and produce a full 3D-rendered image without any  
 126 external API.

### 127 **2.3. Radiative Transfer.**

128 **Primary references:** most of this was adapted from 1D numerical integration  
 129 models I implemented in the MIT course 12.815 (Atmospheric Radiation and Con-  
 130 vention), almost entirely based on the landmark textbook "Principles of Planetary  
 131 Climate" by Raymond Pierrehumbert [1].

132 Radiative transfer models track the heat exchange between a star and planet  
 133 (considering both ground and atmospheric layers) through thermal radiation, typically  
 134 after making several simplifying assumptions. For this project, we have three major  
 135 simplifying assumptions

- 136 1. Everything behaves like a blackbody (a perfect absorber and emitter) with  
 137 a radiation flux governed by the Stefan-Boltzmann law  $F = \sigma T^4$  (given in  
 138 Watts per unit area)
- 139 2. We consider only absorption/reflection of two "wavelengths": shortwave and  
 140 longwave. We assume that all radiation from the star is shortwave, and all  
 141 radiation from the planet is longwave. This is approximately true for the  
 142 Earth and Sun, whose wavelength distributions overlap by less than 1%, and  
 143 a common approach for two-stream equations in atmospheric physics, but  
 144 neglects the complicated absorption profiles of real gases.
- 145 3. The planet is a tidally locked "water world" with a simple CO<sub>2</sub>-like green-  
 146 house atmosphere. These choices were made due to time constraints, since  
 147 too many features would have taken longer to implement and evaluate. This  
 148 allows us to examine the ice-albedo feedback (since the water surface can  
 149 freeze/melt based on temperature), and tidal locking gives a strong tempera-  
 150 ture gradient (which is static at equilibrium, since the same side always faces  
 151 the sun) with a hot side and cool side. With some small modifications, it  
 152 could also be used to study the runaway greenhouse effect (but unfortunately  
 153 I didn't have time to implement it before the deadline).

154 We will revisit these assumptions, and how they can be improved upon, when we  
 155 discuss more detailed models in [subsection 4.1](#) (Future Work). I hope to implement  
 156 many of these features in my research, but my priority for this project was to become  
 157 familiar with ray-traced atmospheric modeling using a relatively simple toy problem  
 158 that tries to capture as much physics as possible (without missing the forest for the  
 159 trees).

#### 160 **2.3.1. Shortwave Radiation.**

161 Shortwave radiation comes from the star, and is primarily composed of near-  
 162 infrared and visible light (as well as small but significant quantities of ultraviolet  
 163 light). We can represent these by casting rays of uniform intensity from a nearby  
 164 disk of radius  $R_{planet}$  representing the stellar direct beam, neglecting effects like limb  
 165 darkening (where the edges of the beam aren't as bright as the center). The total  
 166 incoming energy, which is evenly divided between all the rays, is based on the black-  
 167 body radiation formula  $F = \sigma T^4$  but also depends on the flux at that orbital distance  
 168 (the ratio of the star's surface to the orbit's "surface" area) and the cross-sectional  
 169 area of the planet

$$E = \frac{R_{star}^2}{R_{orbit}^2} \sigma T_{star}^4 \cdot (\pi R_{planet}^2)$$

170 The model assumes that the atmosphere does not interact significantly with short-  
 171 wave radiation, which is a reasonable approximation for a simple atmosphere without  
 172 clouds or ozone. This allows us to consider only the intersection with the planetary  
 173 surface, which may absorb ( $p = 1 - \alpha$ ) or reflect ( $p = \alpha$ ) based on the surface albedo.  
 174 Since this approach requires computing only one intersection, the shortwave step of  
 175 the model is relatively fast (approximately 100 million rays per second).

176 Since albedo is temperature-dependent (ice is more reflective than liquid water)  
 177 it can produce interesting phenomena like ice-albedo feedback, where a planet can  
 178 suddenly shift from an "ice age" state (where less starlight is absorbed) to a warm  
 179 state (where the surface is thawed and absorbs more starlight) depending on its initial  
 180 state, due to hysteresis around unstable equilibrium points.

### 181 2.3.2. Longwave Radiation.

182 Since we assume a pure greenhouse atmosphere, all of the atmospheric layers (as  
 183 well as the ground) can absorb and emit longwave radiation. As before, the emission  
 184 is governed by the blackbody emission formula  $F = \sigma T_{i,j}^4$ , where  $T_{i,j}$  is the local  
 185 temperature at the location of emission.

186 The direction of emission is assumed to be spherically symmetric (as in blackbody  
 187 radiation) so the emitted rays may end up being absorbed by a different location in  
 188 a different layer (possibly the surface) or escaping into space. The probability of  
 189 absorption decays exponentially with optical thickness  $\tau$  (which measures the opacity  
 190 of a path, telling us how likely it is that the gas will absorb the light) by the Beer-  
 191 Lambert-Bouguer law

$$I = I_0 e^{-\tau} \quad \rightarrow \quad p_{transmit} = e^{-\tau}, \quad p_{absorb} = 1 - e^{-\tau}$$

192 Using this, we can more efficiently assign an absorption layer by first drawing  
 193 a probability  $p$ , then using a binary search to find the first layer that crosses that  
 194 threshold (i.e. the minimum "thick enough" distance to be absorbed with this proba-  
 195 bility). This avoids needing to compute a bounce at every layer, although it still adds  
 196 a significant computational expense to each longwave bounce (which takes about 6  
 197 seconds to trace 25 million rays, compared to only 2 seconds for rendering, and only  
 198 a quarter of a second for shortwave).

199 We can keep track of the evolving temperatures with a rank-3 tensor (the pro-  
 200 gramming kind, not the mathematical kind) where each layer (101 by default) is  
 201 broken up into a grid ( $24 \times 16$  cells by default). The cells are sampled uniformly and  
 202 emit a certain amount of longwave energy per time step based on their temperature  
 203 and area (larger cells, near the equator, emit more total energy than smaller cells near  
 204 the poles). This strikes a reasonable balance between simplicity, resolution and phys-  
 205 ical accuracy for this type of toy model, but ideally it would use a weighted sampling  
 206 method that draws more rays from high-emission locations instead of compensating  
 207 by scaling their energies.

### 208 2.3.3. Non-Radiative Heat Transfer.

209 Alternative forms of heat transfer, such as conduction and convection, are ap-  
 210 proximated using a simple weighted Jacobi method. This allows heat to diffuse to  
 211 neighboring cells within the same layer (two vertical, two horizontal) by computing  
 212 the average temperature and using a weighted sum to update the values

$$\tilde{T}_{i,j} = (1 - w) \cdot T_{i,j} + w \cdot \frac{T_{i+1,j} + T_{i-1,j} + T_{i,j+1} + T_{i,j-1}}{4}$$

213 The weight parameter  $w$  controls the extent of diffusion. When  $w = 0$ , there  
 214 is no temperature diffusion at all, and when  $w = 1$  it gives the classical Jacobi  
 215 method (where each cell is simply the average of its 4 nearest neighbors). Intermediate  
 216 values are a weighted sum of these two limits, and low values of  $w$  lead to weakly  
 217 diffusive behavior, which is most consistent with real tidally locked planets where  
 218 we see large temperature gradients (but remains well above absolute zero, even in  
 219 permanent darkness).

#### 220 2.3.4. Physical and Geometric Setup.

221 Unlike the general ray tracer, the planetary ray tracer imposes a very restricted  
 222 geometric setup for performance purposes. The planet and its atmosphere are rep-  
 223 resented by a set of concentric spheres, where the surface (layer 1) has a radius of  
 224  $R_{planet}$  and the radii of all subsequent layers are determined by the pressure profile  
 225 using the pressure-altitude equation (derived from hydrostatic equilibrium)

$$R_j = R_{planet} - H \cdot \log(P_j/P_{surface})$$

226 Effectively, the radius increases in proportion to the log of the pressure ratio (we  
 227 subtract because  $P_j < P_{surface}$  for any  $j > 1$ , so the logarithm is always negative).  
 228 The scale height is a property of the planet that represents the degree of "puffiness"  
 229 of an atmosphere. It is effectively just a scaling factor for this logarithmic relationship,  
 230 which is typically obtained by fitting a pressure profile to an altitude profile (I use  
 231  $H = 8.5 \text{ km}$  by default, which is similar to Earth's atmosphere).

232 A common approach in atmospheric science is to use pressure as the primary ver-  
 233 tical coordinate and calculate all other physical attributes from it, since it is mono-  
 234 tonically decreasing and is more physically meaningful than altitude. In addition  
 235 to the pressure-altitude relation mentioned above, we can also compute the initial  
 236 temperature profile of a given layer using

$$T_j = T_{surface} \cdot (P_j/P_{surface})^\gamma$$

237 Where  $\gamma$  is a scaling parameter between 0 and 1. Setting it to 0.6655 reproduces  
 238 the widely used 1976 US Standard Atmosphere, specifically recreating the tropo-  
 239 spheric profile (the lowest atmospheric layer). The current model does not include  
 240 any atmospheric stratification, since there is no physical mechanism to preserve those  
 241 in the current model (which neglects convection and clouds). Since the purpose of this  
 242 radiative transfer model is to reach thermal equilibrium, the temperature profile is  
 243 not preserved when the model is run, but it provides a useful initial condition to speed  
 244 up convergence and (in cases of hysteresis) potentially leads to different equilibrium  
 245 points.

246 Similarly, we can calculate the optical thickness (opacity) profile using an almost  
 247 identical scaling relationship defined by a parameter  $n$ , which is not bound to the  $[0,1]$   
 248 range. I use  $n = 4$  by default, representing a fairly strong pressure-broadening effect,  
 249 but other values can be physically reasonable as well

$$\tau_j = \tau_{surface} \cdot (P_j/P_{surface})^n$$


250 Since we are using a static pressure profile, the optical thickness is also static (un-  
 251 like temperature). In reality, optical thickness is a complicated property that increases  
 252 with both pressure and temperature. Since it also increases absorption (increasing  
 253 both pressure and temperature), these interdependencies lead to sophisticated feed-  
 254 back loops that remain an active area of research. I hope to include that in future

255 versions of this model, but I didn't want to tackle it this early since it would have  
256 made the scope of this project too broad.

### 257 **3. Results.**

#### 258 **3.1. Renderer.**

259 The following image was generated using two green triangles (representing the  
260 ground) and 20 spheres of various materials and randomly assigned colors (though  
261 some may be out of frame). The colors sampling is somewhat arbitrary, but it was  
262 mostly intended to highlight certain visual effects. For example, glass spheres have a  
263 slight absorption (between 0% and 20% per channel) to make them more interesting  
264 and believable than a (completely unphysical) perfectly refracting sphere.



265 If you look closely, you can see several limitations in this camera-based rendering  
266 technique. Perhaps the most noticeable is that light has no direction at all, and kind  
267 of resembles a very cloudy day where light is scattered uniformly in all directions. The  
268 shadows are also rather unphysical, compared to a situation with a clear and directed  
269 light source, since they are entirely based on bounces (shaded areas occur rays cast  
270 from the camera typically can only "access the sky" after multiple bounces).

271 Additionally, there are a few apparent flaws that are actually correct, but ad-  
272 mittedly still kind of weird. Perhaps the most noticeable is that highly transparent  
273 spheres don't really cast a shadow. Refracting spheres with higher absorption cast a  
274 faint shadow, but glass spheres with little absorption don't cast much of a shadow at  
275 all. This is actually not physically inaccurate, since perfect glass spheres behave in  
276 some pretty unintuitive ways. However, the lack of lensing (or any direct light source  
277 at all) makes them still look rather wrong, but some of the features in radiative  
278 transfer might help with this.

#### 279 **3.2. Physical Phenomena.**

280 While I based many of the physical parameters on the Earth and Sun, this isn't a  
281 particularly realistic model and none of the numbers should be taken seriously. There  
282 is a lot of physics that needed to be ignored here to focus on high-performance GPU  
283 computation on a tight timeline. I hope to develop a physically accurate model of  
284 the future, but the goal here was to develop a ray-traced toy model that can at least  
285 demonstrate some known climate feedbacks.

### 3.2.1. Ice-Albedo Feedback.

286

287

288

289

290

291

292

293

294

295

296

297

298

299

300

The ice-albedo feedback is a phenomenon driven by an inverse relationship between albedo and temperature (i.e. ice is more reflective than water). Essentially, cold planets absorb less light and are difficult to warm, while hot planets absorb more light and are difficult to cool. This results in multiple equilibrium points for certain configurations, and it means that different initial conditions (or temporary changes in conditions) can change which equilibrium point a planet will converge to.

To demonstrate this effect, I used a fairly extreme test case where the planet begins in a frozen state (200  $K$  across the entire surface) and is allowed to reach equilibrium before the star temporarily goes from 6000  $K$  to 9000  $K$  (a completely unphysical test case) to ensure the planet is knocked into another equilibrium state. I'm sure this is greater than necessary, but I didn't have enough time to test the sensitivity thoroughly. Additionally,  $\alpha_{liquid} = 0.1$  and  $\alpha_{ice} = 0.9$ , meaning that a cold surface absorbs only 10% of incident light and a warm one absorbs 90%. These aren't too far off from the real numbers, but a bit on the extreme side.




FIG. 1. The range of temperatures (in Kelvin) and heat exchange (in Joules per timestep) on the planetary surface. The orange line indicates the maximum value (hottest and fastest heating) and the blue line indicates the minimum value (coldest and fastest cooling)

301

302

303

304

305

306

307

308

309

310

311

312

313

314

315

316

317

The temperature curve shows that a planet at 200  $K$  (completely frozen) initially approaches an equilibrium around 210  $K$  (an ice age) but after an extreme temporary rise in temperature, it instead converges around 375  $K$  (warmer than the Earth, but plausible for a planet without clouds near a 6000  $K$  star). This temperature change is driven by the heat transferred  $Q$ , which measures the difference between incoming and outgoing energy for each grid cell.

Looking at  $Q$  (essentially the derivative of  $T$ ), we can see that the sudden warming leads to a period of "overshooting" the equilibrium point. This is likely because the time step did not change, and may have been too coarse for such a major change in conditions. These fluctuations diminish when the star returns to normal, but they still appear to be more variable than the initial phase (starting from a frozen state).

Although the test case is quite extreme, and realistic examples of the ice-albedo feedback exhibit much smaller swings, we can see a clear example of hysteresis due to how the radiative balance changes with temperature. [Disclaimer: due to time constraints, I didn't want to test more realistic values and risk ending up with results without anything physically interesting. I also used only zero diffusion ( $w = 0$ ) because that made hysteresis far more likely]



### 3.3. Performance.

The two "versions" of the ray tracer are built on similar principles, sharing a significant amount of code, even if the overall functionality is quite different. The biggest shared DNA is in how they represent rays, how they handle spherical intersections (the most useful geometry for planetary analysis), and how they handle diffuse reflections/randomly directed emission.

However, the planetary model is more computationally intensive due to the inclusion of thermal absorption/emission and heat transfer, plus an alternate (iterative) method of determining intersections for geometries that only include concentric spheres. As a result, I cut some unnecessary features to maintain a decent thread count and avoid sacrificing too much performance.

All of my results were generated using 25 million rays per bounce, unless otherwise specified, since that fits comfortably within my card's memory (consuming approximately 14 GB in total) and takes approximately 2 seconds to compute a single rendering bounce, or a quarter of a second to compute shortwave bounce. Longwave bounces take a bit longer to complete, typically around 6 seconds for the same number of rays.

#### 3.3.1. Rendering Performance.

To test how performance scales with the number of rays, I ran the renderer using the default settings (2 triangles forming the ground, 20 spheres with randomly assigned surface materials) using a variable number of rays and only a single bounce.


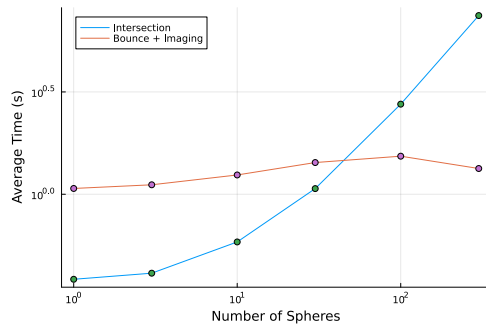


FIG. 2. None

We can see strong linear scaling, where 10× as many rays take approximately 10× as long. There is sometimes some fluctuation when the ray count is really small, though this unusually consistent run suggests that these are random and not caused by overhead. For the most part, the linear scaling seems to hold very well for both stages (computing intersections, and tracking bounces/pixel colors), which take approximately 0.8 and 1.2 seconds respectively for 2.5 million rays.

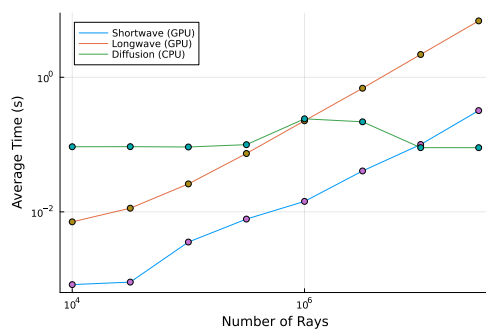
FIG. 3. *None*

345 When we fix the number of rays to 25 million, and instead vary the number of  
 346 spheres, we can see linear scaling in computing the intersections when the number  
 347 of objects is sufficiently large (e.g. 300 spheres take  $10\times$  longer than 30). This is  
 348 expected due to my implementation, which checks every object for a collision (more  
 349 efficient implementations use spatial queries like octrees), but implementing this on  
 350 GPU is complex and well beyond the scope of this project.

351 Unlike the scaling with respect to rays, we see a significant amount of overhead  
 352 when the number of spheres is low. This is in part due to fixed steps that every ray  
 353 must execute (before and after the object loop) and because I don't vary the number  
 354 of triangles (there are always 2, no matter how many spheres there are).

355 Additionally, although the bounce/imaging step depends on the objects, it's in-  
 356 dependent of the number of objects. Since the first intersected object was already  
 357 determined in the previous step, the bounce step just needs to look up the properties  
 358 of a single object (to get materials, color, etc). As a result, it remains more or less  
 359 constant across all the runs, except for minor fluctuations.

360 **3.3.2. Radiative Transfer Performance.** Similar to the rendering case, we  
 361 can see that GPU performance scales roughly linearly with the number of rays once  
 362 there is a sufficient quantity. When the number of rays is small, there are some  
 363 noticeable fluctuations in performance, but they appear to be more or less random  
 364 (rather than a fixed overhead cost).

FIG. 4. *None*

365 The shortwave step is generally quite fast, outperforming either of the steps in  
 366 the rendering case (a quarter of a second for 25 million rays). This is primarily

367 due to the simplifications described in [subsection 2.3.1](#). The longwave step takes  
368 significantly longer, typically around 6 seconds for 25 million rays when using a 101-  
369 layer planet/atmosphere, so I'm not sure if the "simplifications" (like the probabilistic  
370 binary search to determine the absorbing layer) are actually improving performance.  
371 This takes as long to run as the renderer's intersections with 100 spheres, though the  
372 longwave bounce both cuts and adds quite a bit so it's hard to tell (without profiling)  
373 where the performance hit is happening. I'll have to explore this further.

374 Additionally, the diffusion step is currently being handled entirely on CPU. It's  
375 independent of the number of rays, and (unless the number of rays is really small)  
376 has a negligible impact on performance, taking about 0.1 second for the resolution I  
377 used (slightly over 1% of the total runtime for 25 million rays). I plan to parallelize  
378 this step as well, since it would likely run faster and permit higher spatial resolutions,  
379 plus it would significantly reduce data transfer between the CPU and GPU, but it  
380 wasn't worth the effort when the longwave radiation step overwhelmingly dominates  
381 the runtime.

## 382 **4. Discussion.**

### 383 **4.1. Future Work.**

384 I plan to completely redo the radiative transfer form of the ray tracer over the  
385 summer, since this is really just a toy model intended to get my hands dirty with  
386 applied ray tracing in scientific computing. Unlike the renderer (which has a plethora  
387 of useful resources), this implementation was almost entirely self-guided and there is  
388 plenty of room for improvement. As a result, however, it also has a great deal of  
389 research potential (in my opinion) because current implementations are often rather  
390 limited, aimed primarily Earth's climate models rather than exoplanets, and many are  
391 closed-source commercial tools that are not easily accessible to academic researchers.

392 The current model neglects many important physical phenomena, and really just  
393 focuses on a simple model of thermal absorption/emission and albedo feedbacks. Real  
394 atmospheres involve many feedbacks between many of the physical parameters - for ex-  
395 ample, optical thickness/opacity is highly sensitive to both pressure and temperature,  
396 but that isn't being accounted for here. Additionally, the simple longwave/shortwave  
397 breakdown (though common in atmospheric science) is not particularly accurate, and  
398 detailed wavelength-dependent spectra are a crucial component for studying real ex-  
399 oplanets. Finally, the radiative impact of clouds is a complex topic that (at least in  
400 part) would greatly benefit from ray tracing, since they really can't be solved ade-  
401 quately without accounting for their detailed 3D structure (compared to most other  
402 atmospheric problems, which can be solved in 1D without too much loss of informa-  
403 tion).

404 I'm currently working on my thesis proposal, in which I hope to explore all three  
405 of these issues. The first is perhaps the hardest to truly "solve", since many of these  
406 feedback cycles involve significant amounts of uncertainty (especially regarding optical  
407 thickness and the "opacity challenge", which has received significant attention in re-  
408 cent years), so my goal is simply to incorporate the most widely accepted solutions into  
409 a 3D ray tracing environment (most likely something similar to petitRADTRANS).

410 I believe both absorption/emission spectra and cloud interactions represent major  
411 computational challenges, and research opportunities, for GPU ray tracing models.  
412 Traditional methods are quite complex and data-intensive, and typically designed for  
413 serial computing where memory (per thread) is abundant. I'm interested in exploring  
414 the potential for small neural networks to approximate these complicated, uncertain  
415 structures in an efficient, GPU-friendly way.

416 An additional area of improvement (not in research, just in implementation) is  
417 to combine features of both of these methods. A more realistic renderer would first  
418 model light sources and generate texture maps (similar to our temperature grid)  
419 that indicate brightness and color, then use the renderer to capture the image. The  
420 methods I added to the radiative transfer model (which only involves forward ray  
421 tracing) could be combined with the renderer (which currently only involves reverse  
422 ray tracing) to generate 3D images with more realistic lighting, as seen in current  
423 commercial tools and game engines.

424 Finally, I hope to get all of this working in Julia using Nvidia’s Optix library, which  
425 is currently the fastest way to compute ray intersections. It relies on efficient spatial  
426 querying methods and specialized RT cores to compute 100 billion rays per second  
427 (more than 3 orders of magnitude faster than my homebrewed solution). Although  
428 that kind of work is far from my area of expertise, I think it would be a significant  
429 contribution to scientific computing as a whole, since it would make ray tracing far  
430 more accessible to academia and scientific research (because despite using C and C++  
431 for many years, I certainly wouldn’t want to develop libraries or complicated pipelines  
432 within either language).

433 **Acknowledgments.** I’d like to thank Prof. Timothy Cronin (who taught 12.815  
434 Atmospheric Radiation and Convection in parallel to this course) for his helpful ma-  
435 terial on 1D numerical integrals of planetary atmospheres (which I tried my best to  
436 adapt somewhat faithfully into the 3D planetary ray tracing code, within the limited  
437 timeline) and his advice on how to implement some of the finer details.

438

## REFERENCES

- 439 [1] R. PIERREHUMBERT, *Principles of Planetary Climate*, Cambridge University Press, 2010.  
440 [2] P. C. ROBERT AND D. SCHWERI, *GPU-based Ray-Triangle Intersection Testing*, Research Group  
441 on Computational Geometry and Graphics, Institute of Computer Science and Applied  
442 Mathematics, University of Bern, <https://tr.inf.unibe.ch/pdf/iam-04-004.pdf>.  
443 [3] P. SHIRLEY, *Ray tracing in one weekend*. [https://raytracing.github.io/books/](https://raytracing.github.io/books/RayTracingInOneWeekend.html)  
444 [RayTracingInOneWeekend.html](https://raytracing.github.io/books/RayTracingInOneWeekend.html). Accessed: 2023-04.